


GLOBAL TERRORISM

PART 1- GLOBAL TERRORISM ANALYSIS

PART 2- GLOBAL TERRORISM CLASSIFICATION



NAME	ROLL NUMBER
Bhimireddy Sai Gayathri Reddy	CB.EN.U4CSE19012
Harini S	CB.EN.U4CSE19023
Rithika Sri J	CB.EN.U4CSE19025

DATASET LINK : [Global Terrorism Database](#)

DATASET VARIABLES: <https://start.umd.edu/gtd/downloads/Codebook.pdf>

Abstract

The threatened or actual use of illegal force and violence by a non-state actor to attain a political, economic, religious, or social goal through fear, coercion, or intimidation is called Terrorism.

The Global Terrorism Database (GTD) is an open-source database ncluding information on terrorist attacks around the world from 1970 through 2017. The GTD includes systematic data on domestic as well as international terrorist incidents that have occurred during this time period and now includes more than 180,000 attacks. The database is maintained by researchers at the National Consortium for the Study of Terrorism and Responses to Terrorism (START), headquartered at the University of Maryland.

Objective and motivation

In first part our project, we analyze the data available in the global terrorism database (GTD). This analysis will provide insight about terrorism occuring at different parts of the world and their impact on the society.

In the second part of our project, our main aim is to do a binary classification that classifies terror attacks as a 'successful mission' or a 'failed mission' depending upon the independent features provided in the global terrorism database.

REFERENCE

Model Building - https://en.wikipedia.org/wiki/Model_building

Mean Reversion - <https://blog.quantinsti.com/mean-reversion-time-series/>

Time series - <https://builtin.com/data-science/time-series-python>

Feature selection -

<https://machinelearningmastery.com/feature-selection-with-real-and-categorical-data/>

GLOBAL TERRORISM ANALYSIS

GROUP - 4

CB.EN.U4CSE19012 - Gayathri Reddy

CB.EN.U4CSE19023 - Harini S

CB.EN.U4CSE19025 - Rithika Sri J

#Mounting drive

```
from google.colab import drive
drive.mount('/content/drive')
```

Install

```
!pip install pmdarima
```

Collecting pmdarima

Downloading pmdarima-1.8.4-cp37-cp37m-manylinux_2_17_x86_64.manylinux2014_x86_64.manylinux_2_24_x86_64.whl (1.4 MB)

Requirement already satisfied: numpy>=1.19.3 in /usr/local/lib/python3.7/dist-packages (from pmdarima) (1.19.5)

Requirement already satisfied: Cython!=0.29.18,>=0.29 in /usr/local/lib/python3.7/dist-packages (from pmdarima) (0.29.24)

Requirement already satisfied: statsmodels!=0.12.0,>=0.11 in /usr/local/lib/python3.7/dist-packages (from pmdarima) (0.13.1)

Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from pmdarima) (1.1.0)

Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in /usr/local/lib/python3.7/dist-packages (from pmdarima) (57.4.0)

Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.7/dist-packages (from pmdarima) (1.4.1)

Requirement already satisfied: scikit-learn>=0.22 in /usr/local/lib/python3.7/dist-packages (from pmdarima) (1.0.1)

Requirement already satisfied: urllib3 in /usr/local/lib/python3.7/dist-packages (from pmdarima) (1.24.3)

Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.7/dist-packages (from pmdarima) (1.1.5)

Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pmdarima) (2018.9)

Requirement already satisfied: python-dateutil>=2.7.3 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.19->pmdarima) (2.8.2)

Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.7.3->pandas>=0.19->pmdarima) (1.15.0)

```
Requirement already satisfied: threadpoolctl>=2.0.0 in  
/usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.22-  
>pmdarima) (3.0.0)
```

```
Requirement already satisfied: patsy>=0.5.2 in  
/usr/local/lib/python3.7/dist-packages (from statsmodels!  
=0.12.0,>=0.11->pmdarima) (0.5.2)
```

```
Installing collected packages: pmdarima  
Successfully installed pmdarima-1.8.4
```

```
pip install feature-engine
```

```
Requirement already satisfied: feature-engine in  
/usr/local/lib/python3.7/dist-packages (1.1.2)
```

```
Requirement already satisfied: numpy>=1.18.2 in  
/usr/local/lib/python3.7/dist-packages (from feature-engine) (1.19.5)
```

```
Requirement already satisfied: scipy>=1.4.1 in  
/usr/local/lib/python3.7/dist-packages (from feature-engine) (1.4.1)
```

```
Requirement already satisfied: pandas>=1.0.3 in  
/usr/local/lib/python3.7/dist-packages (from feature-engine) (1.1.5)
```

```
Requirement already satisfied: statsmodels>=0.11.1 in  
/usr/local/lib/python3.7/dist-packages (from feature-engine) (0.13.1)
```

```
Requirement already satisfied: scikit-learn>=0.22.2 in  
/usr/local/lib/python3.7/dist-packages (from feature-engine) (1.0.1)
```

```
Requirement already satisfied: python-dateutil>=2.7.3 in  
/usr/local/lib/python3.7/dist-packages (from pandas>=1.0.3->feature-  
engine) (2.8.2)
```

```
Requirement already satisfied: pytz>=2017.2 in  
/usr/local/lib/python3.7/dist-packages (from pandas>=1.0.3->feature-  
engine) (2018.9)
```

```
Requirement already satisfied: six>=1.5 in  
/usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.7.3-  
>pandas>=1.0.3->feature-engine) (1.15.0)
```

```
Requirement already satisfied: threadpoolctl>=2.0.0 in  
/usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.22.2-  
>feature-engine) (3.0.0)
```

```
Requirement already satisfied: joblib>=0.11 in  
/usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.22.2-  
>feature-engine) (1.1.0)
```

```
Requirement already satisfied: patsy>=0.5.2 in  
/usr/local/lib/python3.7/dist-packages (from statsmodels>=0.11.1-  
>feature-engine) (0.5.2)
```

INTRODUCTION

The threatened or actual use of illegal force and violence by a non-state actor to attain a political, economic, religious, or social goal through fear, coercion, or intimidation is called Terrorism.

The main objective of this case study is to analyze the data available in the global terrorism database (GTD). The Global Terrorism Database (GTD) is an open-source database including information on terrorist attacks around the world from 1970 through 2017. The GTD includes systematic data on domestic as well as international terrorist incidents that have occurred during this time period and now includes more than 180,000 attacks.

The database is maintained by researchers at the National Consortium for the Study of Terrorism and Responses to Terrorism (START), headquartered at the University of Maryland.

OVERVIEW:

A detailed analysis on how terrorism spread around the world and the impact caused.

OBJECTIVE:

- To predict future Attacks based on past successful attacks using different models & do model comparison
- To forecast the future trend of Global Terrorism - Number of Attacks using Time Series Analysis

#importing necessary libraries

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import folium
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.seasonal import seasonal_decompose
from datetime import date, timedelta
from pmdarima.arima import auto_arima
from math import sqrt
from sklearn.metrics import mean_squared_error
from feature_engine.creation import CyclicalTransformer
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
```

Understanding the dataset

#loading the dataset

```
GT = pd.read_csv('/content/drive/MyDrive/Global_Terrorism.csv',
encoding = "ISO-8859-1")
```

```
/usr/local/lib/python3.7/dist-packages/IPython/core/
interactiveshell.py:2718: DtypeWarning: Columns
(4,6,31,33,61,62,63,76,79,90,92,94,96,114,115,121) have mixed
types.Specify dtype option on import or set low_memory=False.
interactivity=interactivity, compiler=compiler, result=result)
```

```
#num of columns and rows
```

```
GT.shape
```

```
(181691, 135)
```

```
GT.head()
```

	eventid	iyear	imonth	iday	...	INT_IDEO	INT_MISC	INT_ANY
related								
0	1970000000001	1970	7	2	...	0	0	0
NaN								
1	1970000000002	1970	0	0	...	1	1	1
NaN								
2	1970010000001	1970	1	0	...	-9	1	1
NaN								
3	1970010000002	1970	1	0	...	-9	1	1
NaN								
4	1970010000003	1970	1	0	...	-9	1	1
NaN								

```
[5 rows x 135 columns]
```

```
#Checking on the list of columns
```

```
print(GT.columns.to_list())
```

```
['eventid', 'iyear', 'imonth', 'iday', 'approxdate', 'extended',  
'resolution', 'country', 'country_txt', 'region', 'region_txt',  
'provstate', 'city', 'latitude', 'longitude', 'specificity',  
'vicinity', 'location', 'summary', 'crit1', 'crit2', 'crit3',  
'doubtterr', 'alternative', 'alternative_txt', 'multiple', 'success',  
'suicide', 'attacktype1', 'attacktype1_txt', 'attacktype2',  
'attacktype2_txt', 'attacktype3', 'attacktype3_txt', 'targettype1',  
'targettype1_txt', 'targetsubtype1', 'targetsubtype1_txt', 'corp1',  
'target1', 'natlty1', 'natlty1_txt', 'targettype2', 'targettype2_txt',  
'targetsubtype2', 'targetsubtype2_txt', 'corp2', 'target2', 'natlty2',  
'natlty2_txt', 'targettype3', 'targettype3_txt', 'targetsubtype3',  
'targetsubtype3_txt', 'corp3', 'target3', 'natlty3', 'natlty3_txt',  
'gname', 'gsubname', 'gname2', 'gsubname2', 'gname3', 'gsubname3',  
'motive', 'guncertain1', 'guncertain2', 'guncertain3', 'individual',  
'nperps', 'nperpcap', 'claimed', 'claimmode', 'claimmode_txt',  
'claim2', 'claimmode2', 'claimmode2_txt', 'claim3', 'claimmode3',  
'claimmode3_txt', 'compclaim', 'weaptype1', 'weaptype1_txt',  
'weapsubtype1', 'weapsubtype1_txt', 'weaptype2', 'weaptype2_txt',  
'weapsubtype2', 'weapsubtype2_txt', 'weaptype3', 'weaptype3_txt',  
'weapsubtype3', 'weapsubtype3_txt', 'weaptype4', 'weaptype4_txt',  
'weapsubtype4', 'weapsubtype4_txt', 'weapdetail', 'nkill', 'nkillus',  
'nkillter', 'nwound', 'nwoundus', 'nwoundte', 'property',  
'propextent', 'propextent_txt', 'propvalue', 'propcomment',  
'ishostkid', 'nhostkid', 'nhostkidus', 'nhours', 'ndays', 'divert',  
'kidhijcountry', 'ransom', 'ransomamt', 'ransomamtus', 'ransompaid',  
'ransompaidus', 'ransomnote', 'hostkidoutcome', 'hostkidoutcome_txt']
```

```
'nreleased', 'addnotes', 'scite1', 'scite2', 'scite3', 'dbsource',  
'INT_LOG', 'INT_IDEO', 'INT_MISC', 'INT_ANY', 'related']
```

DATA PREPROCESSING

```
GT = GT.loc[(GT.doubtterr == 0) & (GT.nkill >=0)]  
GT.shape
```

```
(132137, 135)
```

```
#renaming the columns
```

```
GT.rename(columns={'eventid':'ID','iyear':'Year','imonth':'Month','iday':  
'Day','country_txt':'Country','provstate':'state',  
'region_txt':'Region',
```

```
'attacktype1_txt':'AttackType','nkill':'Killed','target1':'Target',  
'nwound':'Wounded','summary':'Summary',  
          'gname':'Group','targettype1_txt':'Target_type',  
'weaptype1_txt':'Weapon_type','motive':'Motive'},inplace=True)
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/frame.py:4308:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame
```

See the caveats in the documentation:

```
https://pandas.pydata.org/pandas-docs/stable/user\_guide/indexing.html#  
returning-a-view-versus-a-copy  
errors=errors,
```

```
#Extracting only important columns
```

```
GT =  
GT[['ID','Year','Month','Day','extended','Country','state','city',  
'Region','latitude','longitude','AttackType','Killed','Wounded','Targe  
t','Group',  
   'success','crit1','crit2','crit3','multiple',  
'Target_type','Weapon_type','vicinity',  
'specificity','suicide','propextent_txt','ishostkid','INT_ANY']]
```

```
#Random 5 rows
```

```
GT.sample(3)
```

	ID	Year	...	ishostkid	INT_ANY
157936	201601200040	2016	...	0.0	0
163164	201606010039	2016	...	0.0	1
147828	201505080012	2015	...	0.0	1

```
[3 rows x 29 columns]
```

```
#num of columns and rows
```

```
GT.shape
```


(132137, 29)

GT.info()

<class 'pandas.core.frame.DataFrame'>

Int64Index: 132137 entries, 0 to 181690

Data columns (total 29 columns):

#	Column	Non-Null Count	Dtype
0	ID	132137 non-null	int64
1	Year	132137 non-null	int64
2	Month	132137 non-null	int64
3	Day	132137 non-null	int64
4	extended	132137 non-null	int64
5	Country	132137 non-null	object
6	state	131934 non-null	object
7	city	131743 non-null	object
8	Region	132137 non-null	object
9	latitude	129568 non-null	float64
10	longitude	129568 non-null	float64
11	AttackType	132137 non-null	object
12	Killed	132137 non-null	float64
13	Wounded	127588 non-null	float64
14	Target	131710 non-null	object
15	Group	132137 non-null	object
16	success	132137 non-null	int64
17	crit1	132137 non-null	int64
18	crit2	132137 non-null	int64
19	crit3	132137 non-null	int64
20	multiple	132136 non-null	float64
21	Target_type	132137 non-null	object
22	Weapon_type	132137 non-null	object
23	vicinity	132137 non-null	int64
24	specificity	132132 non-null	float64
25	suicide	132137 non-null	int64
26	propextent_txt	50322 non-null	object
27	ishostkid	131989 non-null	float64
28	INT_ANY	132137 non-null	int64

dtypes: float64(7), int64(12), object(10)

memory usage: 30.2+ MB

#total null values in each column

GT.isnull().sum()

ID	0
Year	0
Month	0
Day	0
extended	0
Country	0
state	203

city	394
Region	0
latitude	2569
longitude	2569
AttackType	0
Killed	0
Wounded	4549
Target	427
Group	0
success	0
crit1	0
crit2	0
crit3	0
multiple	1
Target_type	0
Weapon_type	0
vicinity	0
specificity	5
suicide	0
propextent_txt	81815
ishostkid	148
INT_ANY	0
dtype:	int64

#Handling Nan in text fields. Assigning Nan as 'Unknown'

GT['Target'].fillna('Unknown', inplace= True)

GT['city'].fillna('Unknown', inplace= True)

GT['state'].fillna('Unknown', inplace= True)

#total null values in each column

GT.isnull().sum()

ID	0
Year	0
Month	0
Day	0
extended	0
Country	0
state	0
city	0
Region	0
latitude	2569
longitude	2569
AttackType	0
Killed	0
Wounded	4549
Target	0
Group	0
success	0
crit1	0
crit2	0

```
crit3            0
multiple         1
Target_type      0
Weapon_type      0
vicinity         0
specificity      5
suicide          0
propextent_txt   81815
ishostkid        148
INT_ANY          0
dtype: int64
```

```
GT = GT.loc[(GT.ishostkid != -9) & (GT.INT_ANY !=-9) ]
```

```
GT=GT.replace('Unknown', np.nan)
GT=GT.replace('Other',np.nan)
```

```
GT.dropna(inplace=True)
GT.shape
```

```
(15691, 29)
```

```
#total null values in each column
```

```
GT.isnull().sum()
```

```
ID            0
Year          0
Month         0
Day           0
extended      0
Country       0
state         0
city          0
Region        0
latitude      0
longitude     0
AttackType    0
Killed        0
Wounded       0
Target        0
Group         0
success       0
crit1         0
crit2         0
crit3         0
multiple      0
Target_type   0
Weapon_type   0
vicinity      0
specificity   0
suicide       0
```

```
propextent_txt      0
ishostkid           0
INT_ANY             0
dtype: int64
```

```
GT.success.value_counts()
```

```
1    15262
0      429
Name: success, dtype: int64
```

```
GT['Day'].value_counts()
```

```
1      628
14     606
21     584
11     576
12     572
2      557
15     554
7      553
10     546
9      541
25     539
13     537
16     532
5      526
19     520
20     515
3      513
27     513
22     485
18     479
4      477
8      476
24     466
26     463
6      453
23     450
17     450
29     446
28     440
30     413
31     239
0        42
```

```
Name: Day, dtype: int64
```

```
...
```

There was data inconsistency in GT['Day']. It had a inconsistent value 0. The actual value must range between 1 - 31. So removing the rows with GT['Day']==0

```

'''
GT.drop(GT.loc[GT.Day==0].index, inplace= True)

#Creating Datetime feature with relevant features
GT['Date'] = pd.to_datetime(GT[['Year','Month','Day']], errors =
'coerce')
#GT.drop(['Day','Month'], axis=1, inplace= True)

#Replacing a large value in GT['Weapon_type']
GT['Weapon_type'].replace({"Vehicle (not to include vehicle-borne
explosives, i.e., car or truck bombs)": "Bombless Vehicle"},inplace
=True)

#Adding new column 'Casualties'
GT['Casualties'] = GT['Wounded'] + GT['Killed']

GT.shape

(15649, 31)

GT.dtypes
ID                                int64
Year                             int64
Month                            int64
Day                              int64
extended                         int64
Country                          object
state                            object
city                             object
Region                           object
latitude                         float64
longitude                       float64
AttackType                       object
Killed                           float64
Wounded                         float64
Target                           object
Group                            object
success                          int64
crit1                            int64
crit2                            int64
crit3                            int64
multiple                        float64
Target_type                      object
Weapon_type                      object
vicinity                        int64
specificity                      float64
suicide                          int64
propextent_txt                   object
ishostkid                       float64
INT_ANY                          int64
Date                            datetime64[ns]

```

```
Casualties          float64
dtype: object
```

```
GT.sample(3)
```

```
          ID  Year  extended  ... INT_ANY      Date
Casualties
6426  197806070003  1978         0  ...      1  1978-06-07
0.0
23733  198411130012  1984         0  ...      0  1984-11-13
0.0
6355  197805200006  1978         0  ...      1  1978-05-20
0.0
```

```
[3 rows x 29 columns]
```

UNDERSTANDING DATA - ANALYSIS

```
print("Country with the most
attacks:",GT['Country'].value_counts().idxmax())
print("Country with the least
attacks:",GT['Country'].value_counts().idxmin())
print("City with the most
attacks:",GT['city'].value_counts().index[0])
print("Region with the most
attacks:",GT['Region'].value_counts().idxmax())
print("Year with the most
attacks:",GT['Year'].value_counts().idxmax())
print("Year with the least
attacks:",GT['Year'].value_counts().idxmin())
print("Group with the most
attacks:",GT['Group'].value_counts().index[1])
print("Most Attack Types:",GT['AttackType'].value_counts().idxmax())
```

```
Country with the most attacks: Iraq
Country with the least attacks: South Yemen
City with the most attacks: Baghdad
Region with the most attacks: Middle East & North Africa
Year with the most attacks: 2014
Year with the least attacks: 1972
Group with the most attacks: Taliban
Most Attack Types: Bombing/Explosion
```

```
#Statistical info on numerical data
```

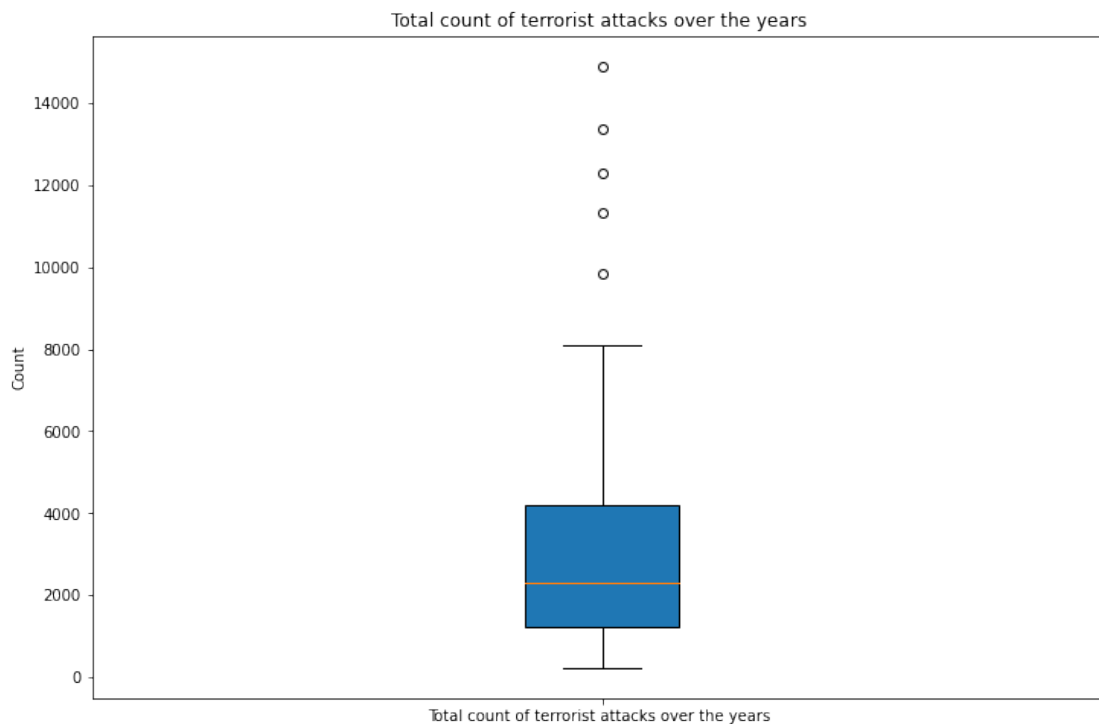
```
GT.describe()[['Killed', 'Wounded', 'Casualties']].round(2)
```

	Killed	Wounded	Casualties
count	160284.00	160284.00	160284.00
mean	2.10	3.20	5.30
std	9.76	36.48	42.58
min	0.00	0.00	0.00

25%	0.00	0.00	0.00
50%	0.00	0.00	1.00
75%	2.00	2.00	4.00
max	1384.00	8191.00	9574.00

INFERENCE FROM DESCRIBE(): Need not rescale 'Killed', 'Wounded' and 'Casualties'.

```
#Box plot for Terrorist attack over the years
year_group = GT.groupby('Year',as_index = False)['ID'].count()
plt.figure(figsize=(12,8))
plt.boxplot(year_group['ID'],patch_artist = True)
plt.title('Total count of terrorist attacks over the years ')
plt.xticks([1],['Total count of terrorist attacks over the years '])
plt.ylabel('Count')
plt.show()
```



```
print('Statistics on total count of terrorist attacks every year
from',year_group['Year'].min(), 'to', year_group['Year'].max(),':')
print('\t Total: ',year_group['ID'].sum())
print('\t Average: ',round(year_group['ID'].mean()))
print('\t Maximum: ',year_group['ID'].max())
print('\t Minimum: ',year_group['ID'].min())
```

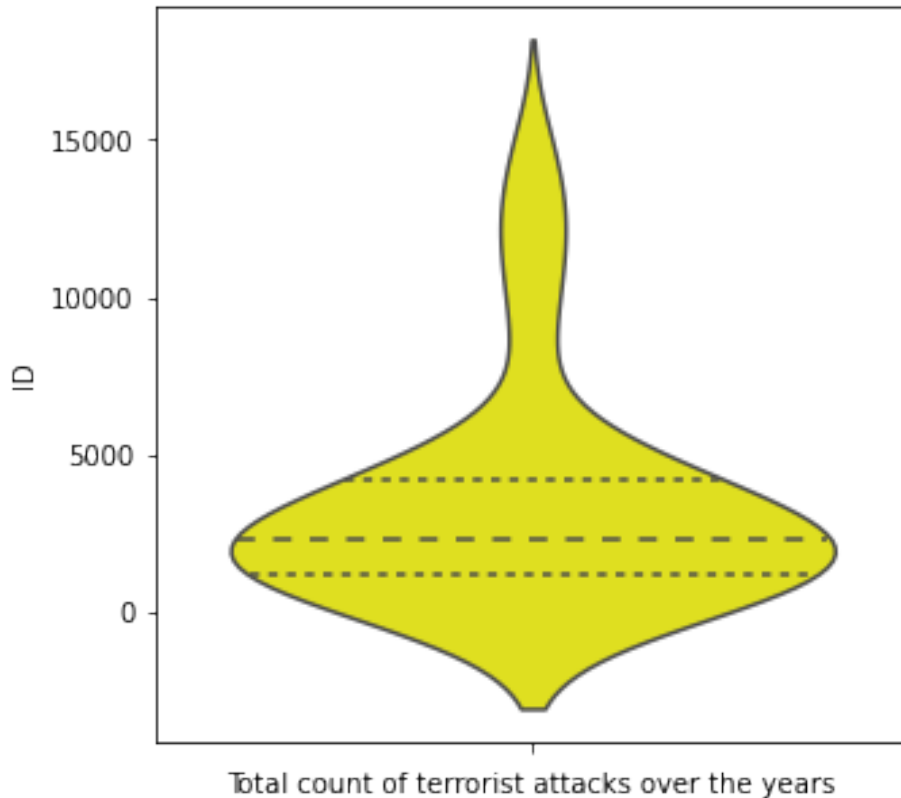
Statistics on total count of terrorist attacks every year from 1970 to 2017 :

```
Total: 160284
Average: 3410
Maximum: 14905
Minimum: 219
```

```

#Violin plot for Terrorist attack over the years
year_group = GT.groupby('Year',as_index = False)['ID'].count()
plt.figure(figsize = (5,5))
plt.xlabel('Total count of terrorist attacks over the years')
sns.violinplot(y= year_group['ID'], inner= 'quartile', color=
'yellow')
plt.show()

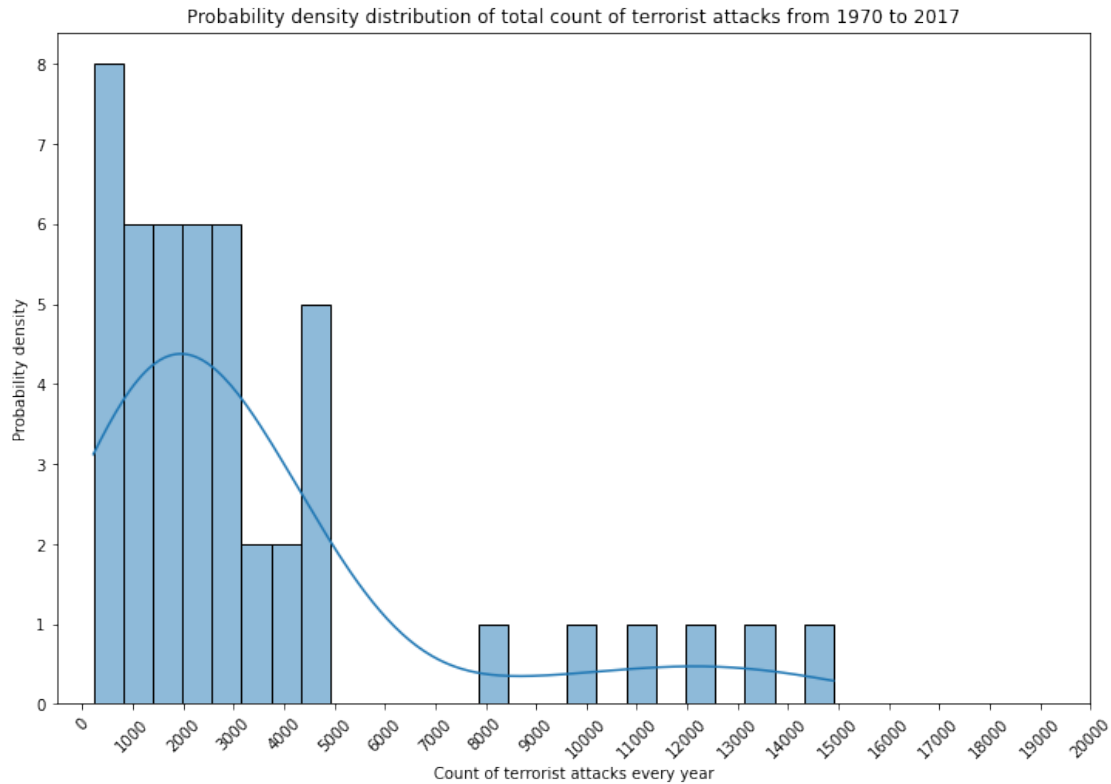
```



```

plt.figure(figsize=(12,8))
ax=sns.histplot(year_group['ID'], bins=25, kde= True)
plt.xlabel('Count of terrorist attacks every year')
plt.ylabel('Probability density')
plt.xticks(range(0,20001,1000),rotation =45)
plt.title('Probability density distribution of total count of
terrorist attacks from 1970 to 2017')
plt.show()

```

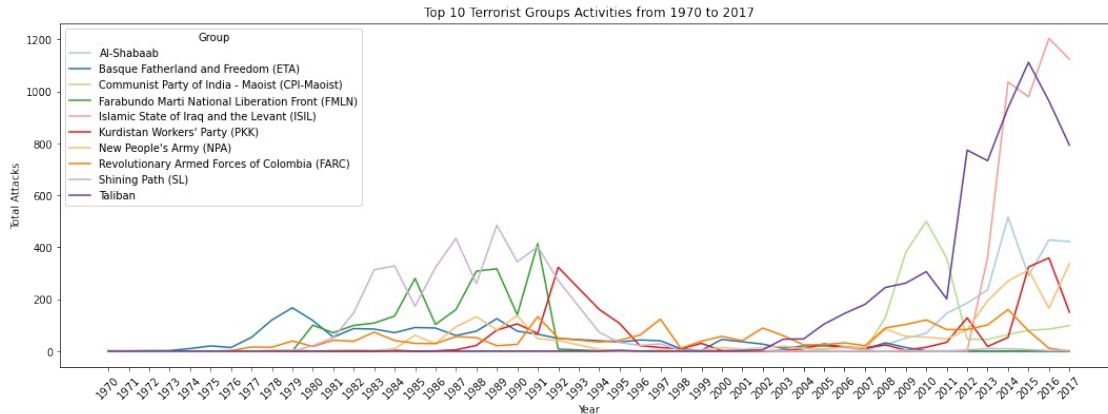



We can conclude that the count of terrorist attack every year is positively skewed.

DATA VISUALIZATION

#Top 10 Terrorist Groups Activities from 1970 to 2017


```
groups_10 = GT[GT.Group.isin(GT.Group.value_counts()[1:11].index)]
pd.crosstab(groups_10.Year,
groups_10.Group).plot(color=sns.color_palette('Paired', 10))
fig=plt.gcf()
fig.set_size_inches(18,6)
plt.xticks(range(1970, 2018, 1), rotation= 45)
plt.ylabel('Total Attacks')
plt.title('Top 10 Terrorist Groups Activities from 1970 to 2017')
plt.show()
```



In the recent years, the Taliban and ISIL are more active than the past decades.

#Count of attacks over the years

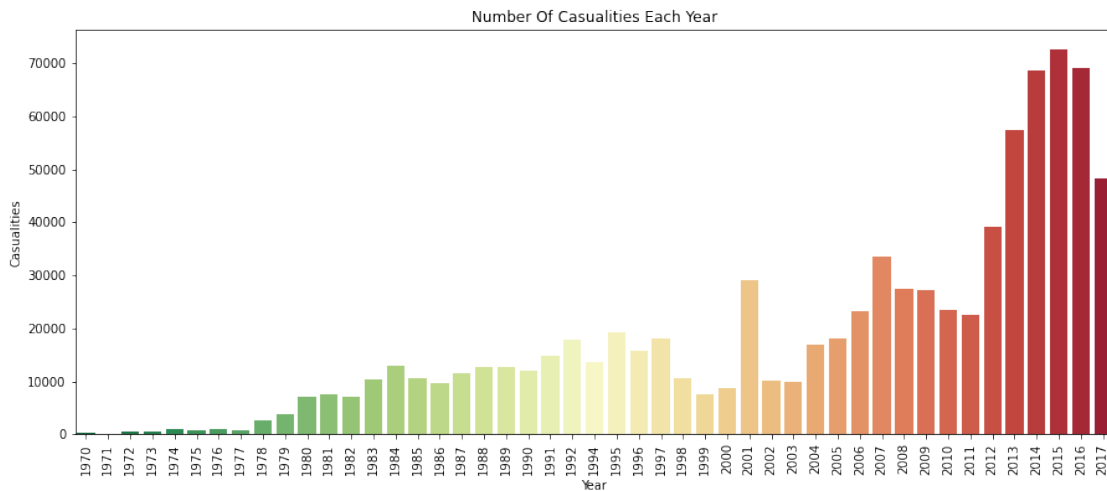
```
x_year = GT['Year'].unique()
y_count_years = GT['Year'].value_counts().sort_index()
plt.subplots(figsize=(15,6))
sns.barplot(x = x_year, y = y_count_years, palette = 'rocket')
plt.xticks(rotation = 45)
plt.xlabel('Attack Year')
plt.ylabel('Number of Attacks each year')
plt.title('Attack_of_Years')
plt.show()
```



As we can see, overall, there has been a rise in the number of terrorist attacks over the years. The highest peak was at 2014 and lowest was at 1972. Seems there is a steady decrease after hitting the highest from 2014 to 2017.

```
plt.subplots(figsize=(15,6))
year_cas =
GT.groupby('Year').Casualties.sum().to_frame().reset_index()
year_cas.columns = ['Year', 'Casualties']
sns.barplot(x=year_cas.Year, y=year_cas.Casualties,
```


```
palette='RdYlGn_r')
plt.xticks(rotation=90)
plt.title('Number Of Casualties Each Year')
plt.show()
```



Inference:

We can see that the number of casualties is high between years 2014 and 2016 and low between the years 1970-1977

```
plt.subplots(figsize=(18,6))
sns.barplot(x= GT['Country'].value_counts()[ :10].index,y=
GT['Country'].value_counts()[ :10].values,palette='inferno')
plt.title('Top Affected Countries')
plt.show()
```

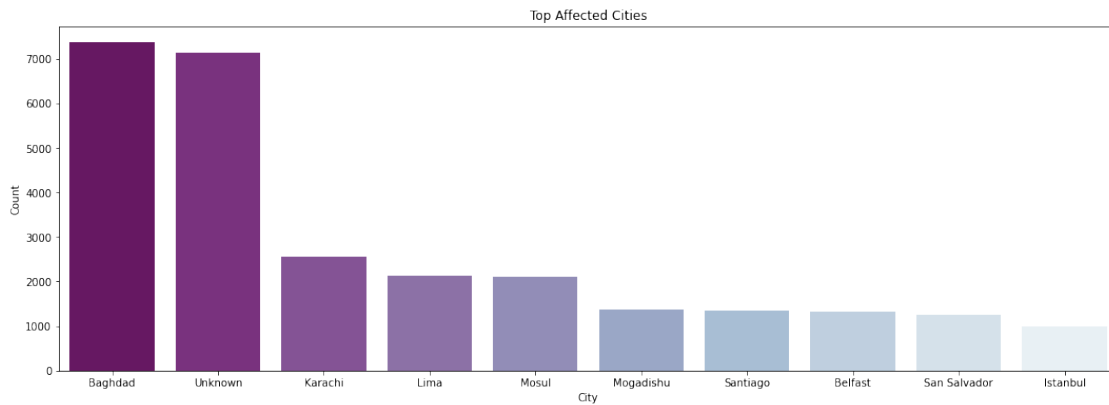


Inference:

The most effected country from Terrorism is Iraq when El Salvador is effected the least.

```
plt.subplots(figsize=(18,6))
sns.barplot(x= GT['city'].value_counts()[ :10].index,y=
GT['city'].value_counts()[ :10].values,palette='BuPu_r')
plt.title('Top Affected Cities')
```

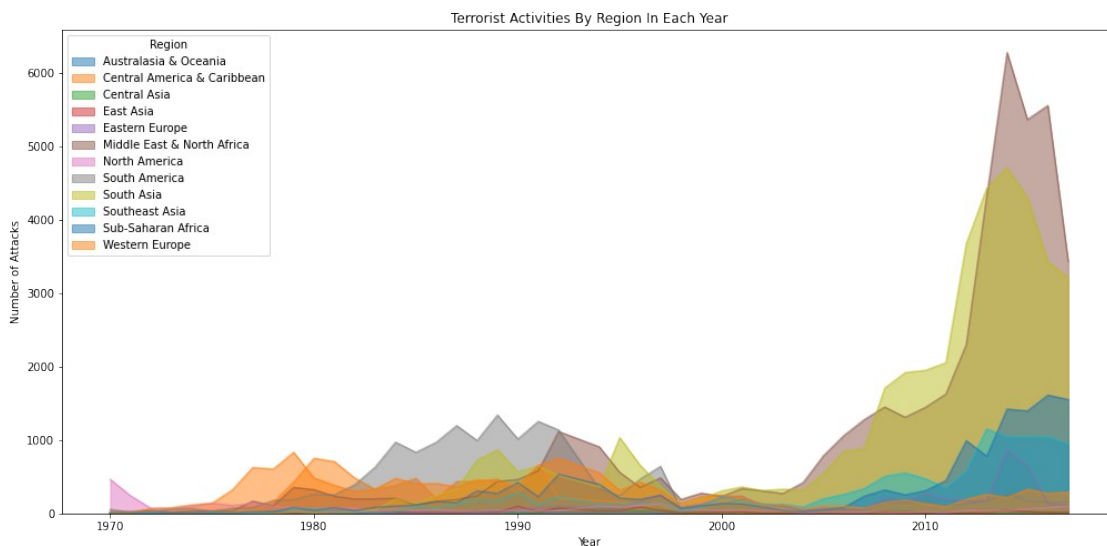
```
plt.xlabel('City')
plt.ylabel('Count')
plt.show()
```



Inference:

The most affected city from terrorism is Baghdad and the least affected city is Istanbul

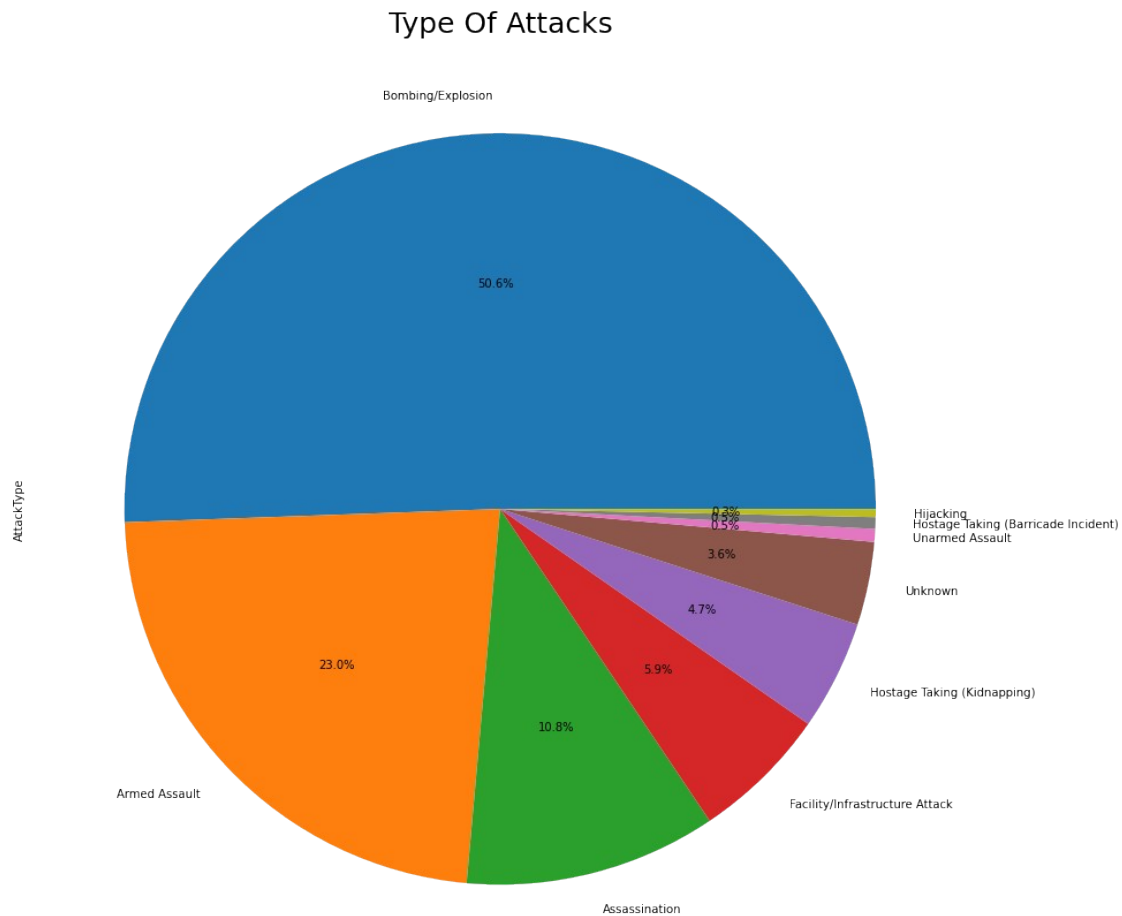
```
pd.crosstab(GT.Year,
GT.Region).plot(kind='area',stacked=False,figsize=(17,8))
plt.title('Terrorist Activities By Region In Each Year')
plt.ylabel('Number of Attacks')
plt.xlabel("Year")
plt.show()
```



Inference:

From The above graph from the year 2010 there is gradual increase of Terrorist activities in the East Asia Region, and also in South Asia Region. In between years 1980-2000 there are high number of Terrorist Activities in South America Region.

```
plt.figure(figsize=(15,15))
GT['AttackType'].value_counts().plot.pie(autopct="%1.1f%%")
plt.title('Type Of Attacks', fontsize=25)
plt.show()
```




Inference:

Bombing or explosion is the most often used for attacking. Where Armed Assault is the second. Hijacking, Hostage Taking and Unarmed Assault are the least used type of Attacks.


```
plt.figure(figsize=(15,15))
GT['Target_type'].value_counts().plot.pie(autopct="%1.1f%%")
plt.title('Target Of Attack', fontsize=25)
plt.show()
```

Target Of Attack



```
plt.figure(figsize=(15,15))
GT['Weapon_type'].value_counts().plot.pie(autopct="%1.1f%%")
plt.title('Weapons Used for Attack',fontsize=25)
plt.show()
```

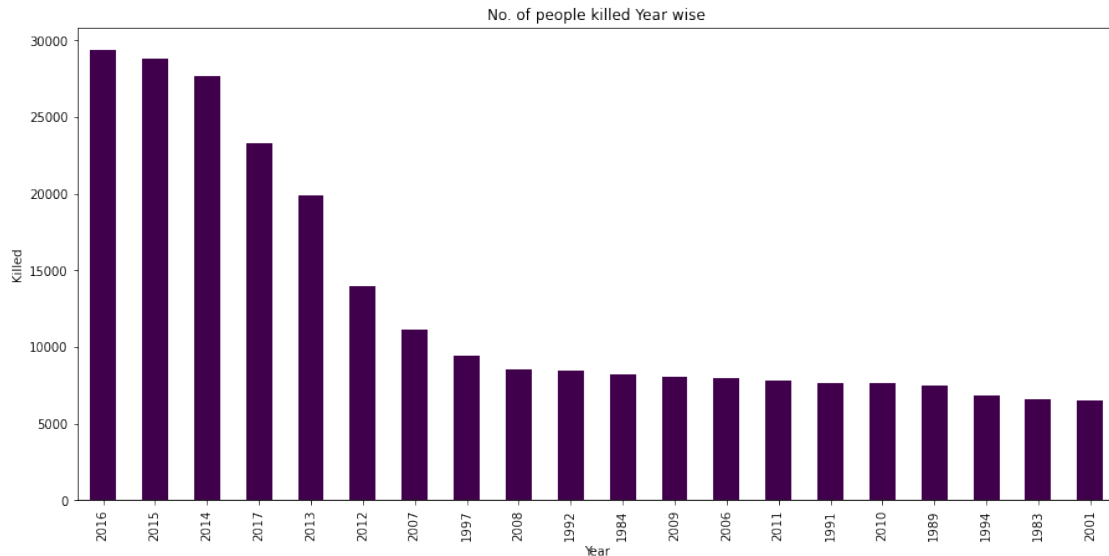
Weapons Used for Attack



Inference:

For the majority of Attacks Explosives are used as weapons and Firearms are in second place.

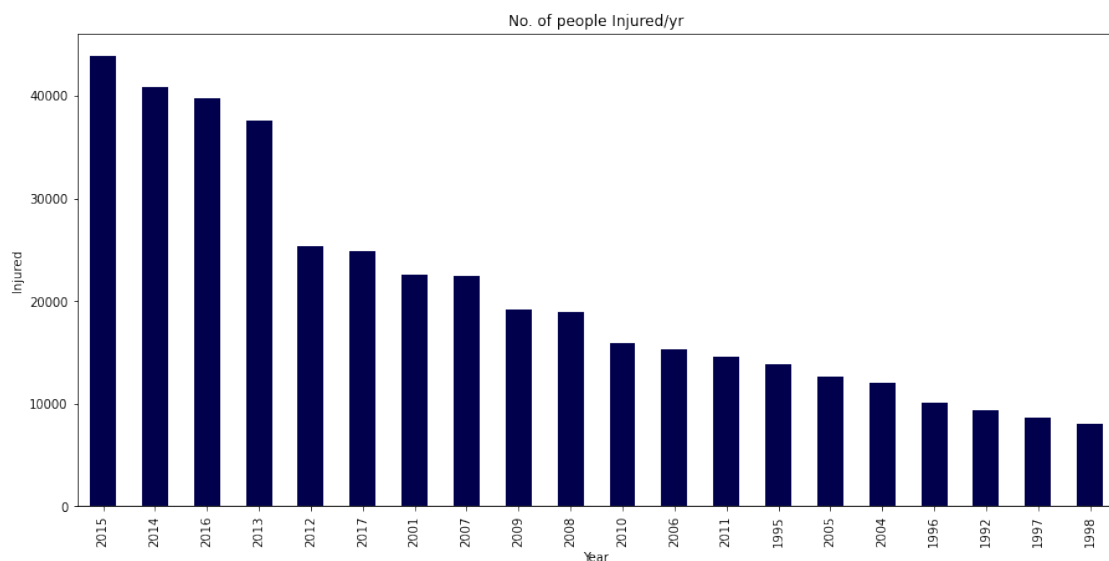
```
plt.figure(figsize = (15,7))
GT.groupby(['Year'])['Killed'].sum().sort_values(ascending =
False).head(20).plot(kind = 'bar', colormap = 'PRGn')
plt.xticks(rotation=90)
plt.title('No. of people killed Year wise')
plt.ylabel("Killed")
plt.show()
```



Inference:

In the year 2016 major number of people are killed, and from 2012 there was a drastic increase of deaths.


```
plt.figure(figsize = (15,7))
GT.groupby(['Year'])['Wounded'].sum().sort_values(ascending =
False).head(20).plot(kind = 'bar', colormap = 'seismic')
plt.xticks(rotation=90)
plt.title('No. of people Injured/yr')
plt.ylabel("Injured")
plt.show()
```



```
plt.figure(figsize = (15,7))
GT.groupby(['Country'])['Killed'].sum().sort_values(ascending =
False).head(10).plot(kind = 'bar', colormap = 'PRGn')
```




```
plt.title('No. of people killed Country wise')
plt.ylabel("Killed")
plt.show()
```



```
print(GT['Country'].unique().shape)
print(GT['Region'].unique().shape)
```

```
(202,)
(12,)
```

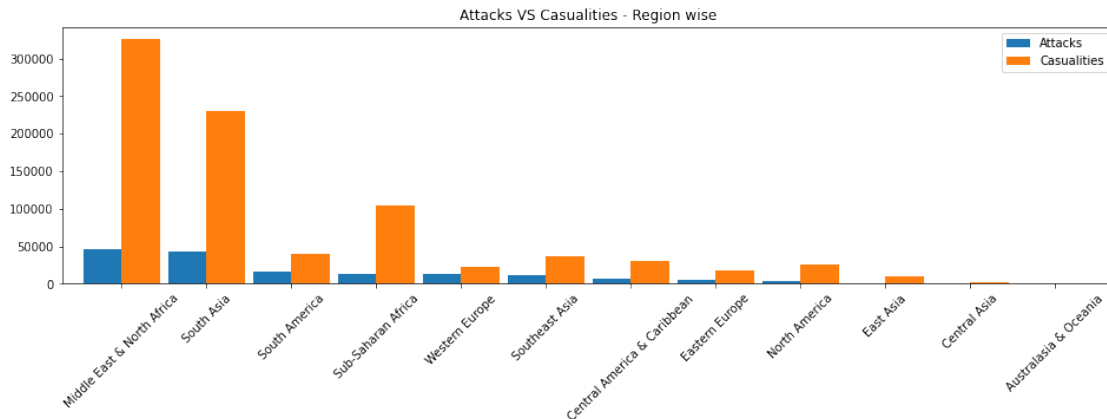
```
count_terror = GT['Country'].value_counts()[ :15].to_frame()
count_terror.columns=['Attacks']
count_kill=GT.groupby ('Country')['Casualties'].sum().to_frame()
count_terror.merge(count_kill,left_index = True,right_index
=True,how='left').plot.bar(width=0.9)
plt.xticks(rotation=45)
fig=plt.gcf()
plt.title("Attacks VS Casualties - Country wise")
fig.set_size_inches(16,4)
plt.show()
```



```

count_terror = GT['Region'].value_counts()[15].to_frame()
count_terror.columns=['Attacks']
count_kill=GT.groupby ('Region')['Casualties'].sum().to_frame()
count_terror.merge(count_kill,left_index = True,right_index
=True,how='left').plot.bar(width=0.9)
fig=plt.gcf()
plt.xticks(rotation=45)
plt.title("Attacks VS Casualties - Region wise")
fig.set_size_inches(16,4)
plt.show()

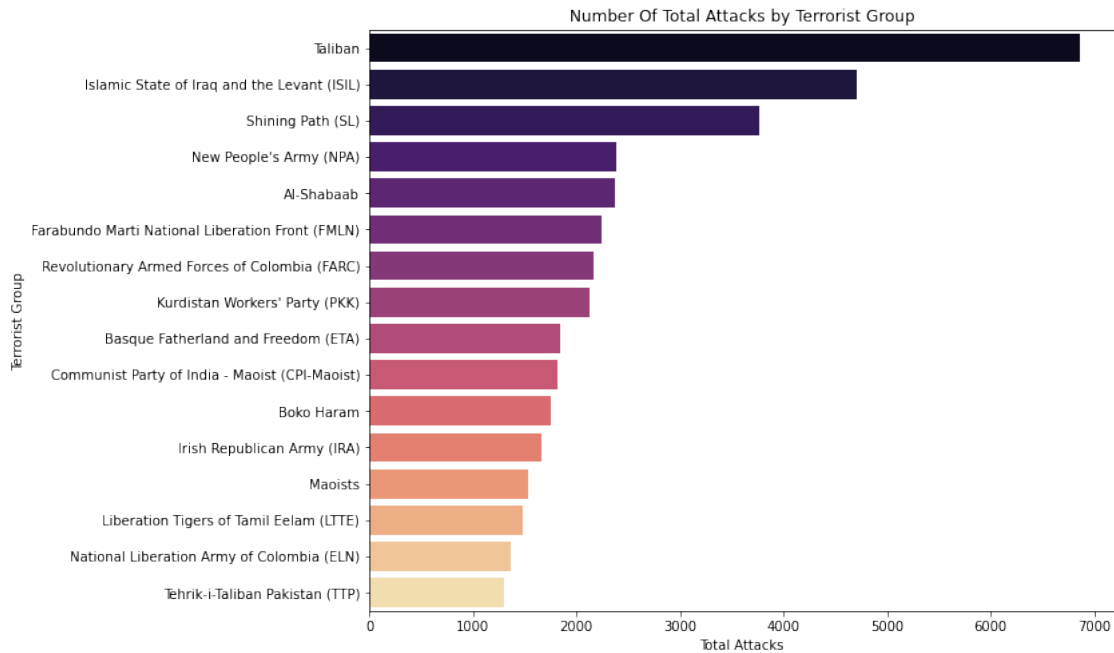
```



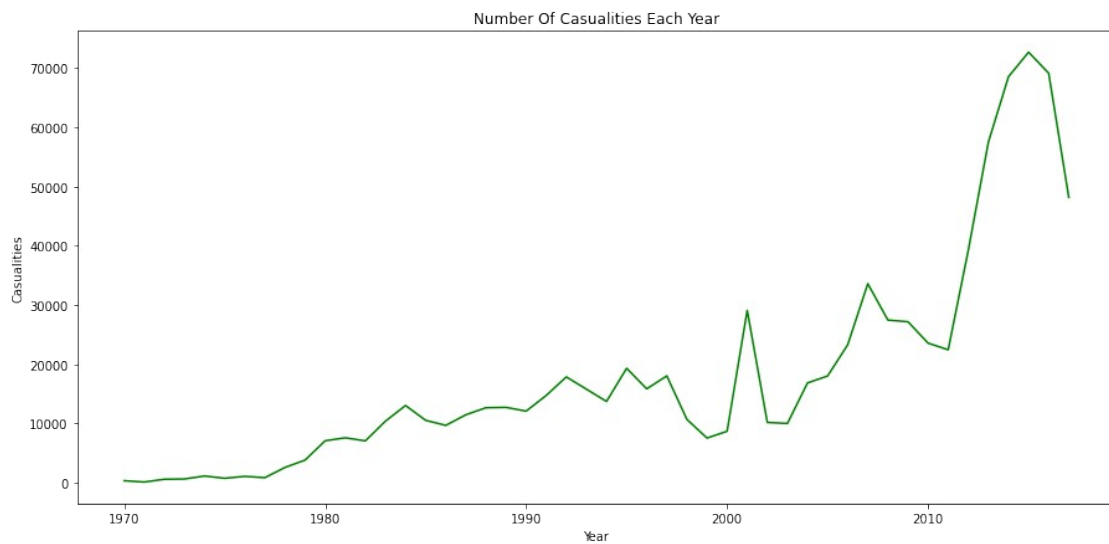
```

group_attacks =
GT.Group.value_counts().to_frame().drop('Unknown').reset_index()[16]
group_attacks.columns = ['Terrorist Group', 'Total Attacks']
plt.subplots(figsize=(10,8))
sns.barplot(y= group_attacks['Terrorist Group'], x=
group_attacks['Total Attacks'], palette='magma')
plt.title('Number Of Total Attacks by Terrorist Group')
plt.show()

```



```
plt.subplots(figsize=(15,7))
year_casual =
GT.groupby('Year').Casualties.sum().to_frame().reset_index()
year_casual.columns = ['Year', 'Casualties']
plt.title('Number Of Casualties Each Year')
sns.lineplot(x='Year', y='Casualties', data=year_casual, color="g")
<matplotlib.axes._subplots.AxesSubplot at 0x7f5ba1f7fc10>
```



=>Year and Casualties are *independent*

```
loca = GT[['latitude', 'longitude']][:8000]
coun = GT['Country'][:8000]
cit = GT['city'][:8000]
```

```

kill = GT['Killed'][:8000]
wound = GT['Wounded'][:8000]

def color(x):
    if x>=30:
        color='red'
    elif ((x>0 and x<30)):
        color='blue'
    else:
        color='orange'
    return color
def size(x):
    if (x>30 and x<100):
        size=2
    elif (x>=100 and x<500):
        size=8
    elif x>=500:
        size=16
    else:
        size=0.5
    return size

map = folium.Map(location=[30,0],tiles='cartodbpositron',zoom_start=2)
for point in loca.index:
    info='<b>Country: </b>'+str(coun[point])+'<br><b>City: </b>:'
    '+str(cit[point])+'<br><b>Killed </b>: '+str(kill[point])
    +'<br><b>Wounded</b> : '+str(wound[point])
    iframe = folium.IFrame(html=info, width=200, height=200)

    folium.CircleMarker(list(loca.loc[point].values),popup=folium.Popup(iframe),radius=size(kill[point]),color=color(kill[point])).add_to(map)

map
<folium.folium.Map at 0x7f5ba1b89c10>

```

INFERENCE:

The map shows that South West Asia has had many attacks compared to the rest of the world


#India - Casualties over Years

```

GTindia=GT[GT['Country']=='India']
year = GTindia["Year"].value_counts()
Year_Attack, Counts_attack = list(year.index), list(year.values)
plt.subplots(figsize=(20,9))
sns.pointplot(x=Year_Attack,y=Counts_attack,color="Red")
plt.xlabel("Year")
plt.xticks(rotation=100)
plt.ylabel("Casualties")

```

```
plt.title("Casualties vs Year - India")
plt.grid(color='b', linestyle='-', linewidth=0.2)
```



INFERENCE:

From the above graph we can see that the Casualties caused by Terrorism in India, same as Global Terrorism, has increased comparatively. The rate of increase in Casualties is not proportional to Years.

```
location_ind=GTindia[['latitude','longitude']][:5000]
city_ind=GTindia['city'][:5000]
killed_ind=GTindia['Killed'][:5000]
wound_ind=GTindia['Wounded'][:5000]
```

```
map1 = folium.Map(location=[20,
78],tiles='cartodbpositron',zoom_start=4)
for point in location_ind.index:
```

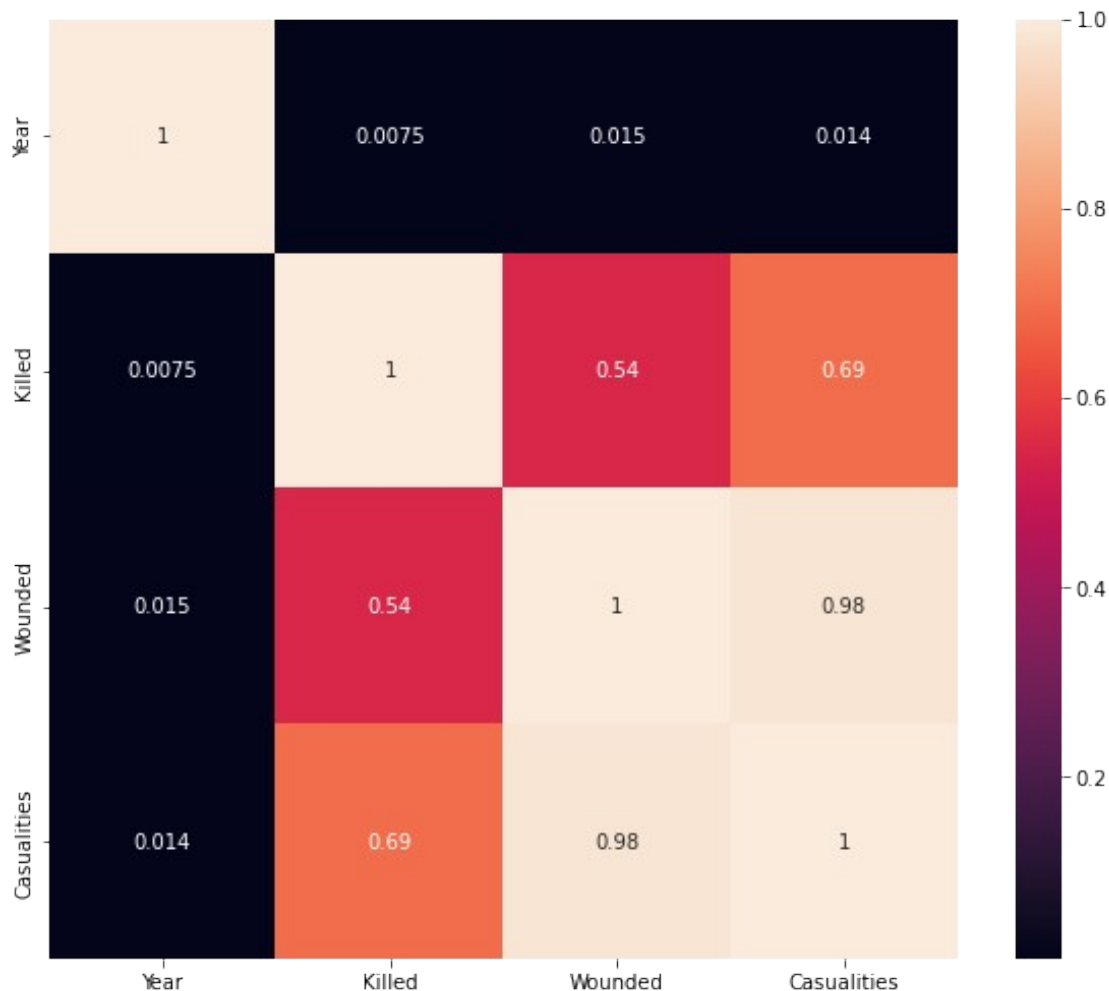
```
folium.CircleMarker(list(location_ind.loc[point].values),popup='<b>City: </b>'+str(city_ind[point])+'<br><b>Killed: </b>'+str(killed_ind[point])+\'
                        \'<br><b>Injured: </b>'+str(wound_ind[point]),radius=size(killed_ind[point]),color=color
(killed_ind[point]),fill_color=color(killed_ind[point])).add_to(map1)
map1
```

```
<folium.folium.Map at 0x7f5b9b3f7c10>
```

INFERENCE

- It is clear from the map that the places that border neighbouring countries suffer more attacks.
- Major places and capitals like Mumbai, Delhi, Gandhinagar, Chennai and Hyderabad has had more casualties (Targetted Cities and States)

```
#Correlation representation
hm = GT[['Year', 'Killed', 'Wounded', 'Casualties']]
plt.figure(figsize=(10,8.5))
sns.heatmap(hm.corr(), annot= True)
plt.show()
```



INFERENCE:

The features Killed - Wounded - Casualty are directly propotional

GLOBAL TERRORISM - BINARY CLASSIFICATION

The main objective of this case study is to classify the data available in the global terrorism database (GTD) as 'successful' or 'failure'. This is a binary classification problem.

```
gt = pd.read_csv('/content/drive/MyDrive/Global_Terrorism.csv',
encoding = "ISO-8859-1")
```

```

/usr/local/lib/python3.7/dist-packages/IPython/core/
interactiveshell.py:2718: DtypeWarning: Columns
(4,6,31,33,61,62,63,76,79,90,92,94,96,114,115,121) have mixed
types.Specify dtype option on import or set low_memory=False.
    interactivity=interactivity, compiler=compiler, result=result)

```

```

gt = gt.loc[(gt.doubtterr == 0) & (gt.nkill >=0)]
gt.shape

```

```

(132137, 135)

```

```

gt

```

	eventid	iyear	imonth	iday	...	INT_IDEO	INT_MISC
INT_ANY	related						
0	1970000000001	1970	7	2	...	0	0
0	NaN						
1	1970000000002	1970	0	0	...	1	1
1	NaN						
2	1970010000001	1970	1	0	...	-9	1
1	NaN						
5	1970010100002	1970	1	1	...	-9	0
-9	NaN						
6	1970010200001	1970	1	2	...	0	0
0	NaN						
...
...
181684	201712310019	2017	12	31	...	0	0
0	NaN						
181685	201712310020	2017	12	31	...	-9	0
-9	NaN						
181688	201712310030	2017	12	31	...	0	0
0	NaN						
181689	201712310031	2017	12	31	...	-9	0
-9	NaN						
181690	201712310032	2017	12	31	...	-9	0
-9	NaN						

```

[132137 rows x 135 columns]

```

```

GT =

```

```

gt[['iyear', 'imonth', 'extended', 'country_txt', 'region_txt', 'nkill', 'nw
ound', 'success', 'crit1', 'crit2', 'crit3', 'multiple',
    'vicinity',
    'specificity', 'suicide', 'attacktype1_txt', 'weaptype1_txt', 'targettype1_t
xt', 'gname', 'propextent_txt', 'ishostkid', 'INT_ANY']]

```

```

GT

```

	iyear	imonth	...	ishostkid	INT_ANY
0	1970	7	...	0.0	0
1	1970	0	...	1.0	1

2	1970	1	...	0.0	1
5	1970	1	...	0.0	-9
6	1970	1	...	0.0	0
...
181684	2017	12	...	1.0	0
181685	2017	12	...	0.0	-9
181688	2017	12	...	0.0	0
181689	2017	12	...	0.0	-9
181690	2017	12	...	0.0	-9

[132137 rows x 22 columns]

```
GT = GT.loc[(GT.ishostkid != -9) & (GT.INT_ANY != -9) ]
GT.shape
```

(70564, 22)

```
GT=GT.replace('Unknown', np.nan)
```

```
GT=GT.replace('Other', np.nan)
```

```
GT.isnull().sum()
```

iyear	0
imonth	0
extended	0
country_txt	0
region_txt	0
nkill	0
nwound	3242
success	0
crit1	0
crit2	0
crit3	0
multiple	0
vicinity	0
specificity	3
suicide	0
attacktype1_txt	3005
weaptype1_txt	6466
targettype1_txt	2183
gname	4494
propextent_txt	51435
ishostkid	119
INT_ANY	0
dtype: int64	

GT

	iyear	imonth	...	ishostkid	INT_ANY
0	1970	7	...	0.0	0
1	1970	0	...	1.0	1

2	1970	1	...	0.0	1
6	1970	1	...	0.0	0
8	1970	1	...	0.0	0
...
181677	2017	12	...	0.0	0
181681	2017	12	...	0.0	0
181683	2017	12	...	0.0	0
181684	2017	12	...	1.0	0
181688	2017	12	...	0.0	0

[70564 rows x 22 columns]

GT.dropna(inplace=True)

GT.shape

(16777, 22)

GT.success.value_counts()

1 16323

0 454

Name: success, dtype: int64

GT.dtypes

iyear	int64
imonth	int64
extended	int64
country_txt	object
region_txt	object
nkill	float64
nwound	float64
success	int64
crit1	int64
crit2	int64
crit3	int64
multiple	float64
vicinity	int64
specificity	float64
suicide	int64
attacktype1_txt	object
weaptype1_txt	object
targettype1_txt	object
gname	object
propextent_txt	object
ishostkid	float64
INT_ANY	int64
dtype:	object

pip install feature-engine

```

Collecting feature-engine
  Downloading feature_engine-1.1.2-py2.py3-none-any.whl (180 kB)
odols>=0.11.1
  Downloading statsmodels-0.13.1-cp37-cp37m-
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (9.8 MB)
ent already satisfied: pandas>=1.0.3 in /usr/local/lib/python3.7/dist-
packages (from feature-engine) (1.1.5)
Requirement already satisfied: scipy>=1.4.1 in
/usr/local/lib/python3.7/dist-packages (from feature-engine) (1.4.1)
Requirement already satisfied: scikit-learn>=0.22.2 in
/usr/local/lib/python3.7/dist-packages (from feature-engine) (1.0.1)
Requirement already satisfied: numpy>=1.18.2 in
/usr/local/lib/python3.7/dist-packages (from feature-engine) (1.19.5)
Requirement already satisfied: pytz>=2017.2 in
/usr/local/lib/python3.7/dist-packages (from pandas>=1.0.3->feature-
engine) (2018.9)
Requirement already satisfied: python-dateutil>=2.7.3 in
/usr/local/lib/python3.7/dist-packages (from pandas>=1.0.3->feature-
engine) (2.8.2)
Requirement already satisfied: six>=1.5 in
/usr/local/lib/python3.7/dist-packages (from python-dateutil>=2.7.3-
>pandas>=1.0.3->feature-engine) (1.15.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.22.2-
>feature-engine) (3.0.0)
Requirement already satisfied: joblib>=0.11 in
/usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.22.2-
>feature-engine) (1.1.0)
Requirement already satisfied: patsy>=0.5.2 in
/usr/local/lib/python3.7/dist-packages (from statsmodels>=0.11.1-
>feature-engine) (0.5.2)
Installing collected packages: statsmodels, feature-engine
  Attempting uninstall: statsmodels
    Found existing installation: statsmodels 0.10.2
    Uninstalling statsmodels-0.10.2:
      Successfully uninstalled statsmodels-0.10.2
Successfully installed feature-engine-1.1.2 statsmodels-0.13.1

```

```

{"pip_warning":{"packages":["statsmodels"]}}

```

```

from feature_engine.creation import CyclicalTransformer

```

```

df = GT[['iyear', 'imonth']]
cyclical = CyclicalTransformer(variables=None, drop_original=True)
X = cyclical.fit_transform(df)
X = X.groupby(X.columns,axis=1).apply(lambda x: round(x,3))

GT = pd.concat([X, GT], axis=1)
GT= GT.drop('iyear', axis=1)

```

```
GT= GT.drop('imonth', axis=1)
GT
```

	iyear_sin	iyear_cos	...	ishostkid	INT_ANY
8	-0.146	0.989	...	0.0	0
9	-0.146	0.989	...	0.0	0
21	-0.146	0.989	...	0.0	0
55	-0.146	0.989	...	0.0	1
56	-0.146	0.989	...	0.0	1
...
181659	-0.000	1.000	...	0.0	1
181665	-0.000	1.000	...	0.0	0
181676	-0.000	1.000	...	0.0	0
181677	-0.000	1.000	...	0.0	0
181681	-0.000	1.000	...	0.0	0

[16777 rows x 24 columns]

```
#GT= pd.get_dummies(data=GT,
columns=['extended', 'country_txt', 'region_txt', 'crit1', 'crit2', 'crit3',
'multiple',
# 'vicinity',
'specificity', 'attacktype1_txt', 'weaptype1_txt', 'targettype1_txt', 'prope
xtent_txt', 'gname', 'ishostkid', 'INT_ANY'])
```

```
from sklearn import preprocessing
# creating instance of labelencoder
labelencoder = preprocessing.LabelEncoder()
# Assigning numerical values and storing in another column
GT['country_txt']= labelencoder.fit_transform(GT['country_txt'])
GT['region_txt']= labelencoder.fit_transform(GT['region_txt'])
GT['attacktype1_txt']=
labelencoder.fit_transform(GT['attacktype1_txt'])
GT['weaptype1_txt']= labelencoder.fit_transform(GT['weaptype1_txt'])
GT['targettype1_txt']= labelencoder.fit_transform(GT['targettype1_txt'])
GT['propxtent_txt']= labelencoder.fit_transform(GT['propxtent_txt'])
GT['gname']= labelencoder.fit_transform(GT['gname'])
GT
```

	iyear_sin	iyear_cos	imonth_sin	...	propxtent_txt
ishostkid	INT_ANY				
8	-0.146	0.989	0.500	...	2
0.0	0				
9	-0.146	0.989	0.500	...	2
0.0	0				
21	-0.146	0.989	0.500	...	2
0.0	0				
55	-0.146	0.989	0.866	...	2
0.0	1				
56	-0.146	0.989	0.866	...	2
0.0	1				

```

...      ...      ...      ...      ...      .
..
181659    -0.000      1.000      -0.000    ...      2
0.0        1
181665    -0.000      1.000      -0.000    ...      2
0.0        0
181676    -0.000      1.000      -0.000    ...      2
0.0        0
181677    -0.000      1.000      -0.000    ...      2
0.0        0
181681    -0.000      1.000      -0.000    ...      2
0.0        0

```

[16777 rows x 24 columns]

GT.shape

(16777, 24)

GT

```

      iyear_sin  iyear_cos  imonth_sin  ...  propextent_txt
ishostkid  INT_ANY
8          -0.146      0.989      0.500    ...      2
0.0         0
9          -0.146      0.989      0.500    ...      2
0.0         0
21         -0.146      0.989      0.500    ...      2
0.0         0
55         -0.146      0.989      0.866    ...      2
0.0         1
56         -0.146      0.989      0.866    ...      2
0.0         1
...      ...      ...      ...      ...      .
..
181659    -0.000      1.000      -0.000    ...      2
0.0        1
181665    -0.000      1.000      -0.000    ...      2
0.0        0
181676    -0.000      1.000      -0.000    ...      2
0.0        0
181677    -0.000      1.000      -0.000    ...      2
0.0        0
181681    -0.000      1.000      -0.000    ...      2
0.0        0

```

[16777 rows x 24 columns]

GT1=

```

GT[['iyear_sin','iyear_cos','imonth_sin','imonth_cos','nkill','nwound'
]]

```

```
GT1['casualties'] = GT1['nkill'] + GT1['nwound']
GT1
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:2:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

	iyear_sin	iyear_cos	imonth_sin	imonth_cos	nkill	nwound
casualties						
8	-0.146	0.989	0.500	0.866	0.0	0.0
0.0						
9	-0.146	0.989	0.500	0.866	0.0	0.0
0.0						
21	-0.146	0.989	0.500	0.866	0.0	0.0
0.0						
55	-0.146	0.989	0.866	0.500	0.0	0.0
0.0						
56	-0.146	0.989	0.866	0.500	0.0	0.0
0.0						
...
...						
181659	-0.000	1.000	-0.000	1.000	0.0	0.0
0.0						
181665	-0.000	1.000	-0.000	1.000	0.0	0.0
0.0						
181676	-0.000	1.000	-0.000	1.000	5.0	0.0
5.0						
181677	-0.000	1.000	-0.000	1.000	0.0	0.0
0.0						
181681	-0.000	1.000	-0.000	1.000	1.0	5.0
6.0						

[16777 rows x 7 columns]

```
GT=
GT.drop(['iyear_sin', 'iyear_cos', 'imonth_sin', 'imonth_cos', 'nkill', 'nwound'], axis=1)
GT = pd.concat([GT1, GT], axis=1)
GT
```

	iyear_sin	iyear_cos	imonth_sin	...	propextent_txt
ishostkid	INT_ANY				
8	-0.146	0.989	0.500	...	2
0.0	0				
9	-0.146	0.989	0.500	...	2

0.0	0					
21	-0.146	0.989	0.500	...		2
0.0	0					
55	-0.146	0.989	0.866	...		2
0.0	1					
56	-0.146	0.989	0.866	...		2
0.0	1					
...
...	...					
181659	-0.000	1.000	-0.000	...		2
0.0	1					
181665	-0.000	1.000	-0.000	...		2
0.0	0					
181676	-0.000	1.000	-0.000	...		2
0.0	0					
181677	-0.000	1.000	-0.000	...		2
0.0	0					
181681	-0.000	1.000	-0.000	...		2
0.0	0					

[16777 rows x 25 columns]

GT.reset_index(drop=True, inplace=True)

GT

	iyear_sin	iyear_cos	imonth_sin	...	propextent_txt
ishostkid INT_ANY					
0	-0.146	0.989	0.500	...	2
0.0	0				
1	-0.146	0.989	0.500	...	2
0.0	0				
2	-0.146	0.989	0.500	...	2
0.0	0				
3	-0.146	0.989	0.866	...	2
0.0	1				
4	-0.146	0.989	0.866	...	2
0.0	1				
...
.	...				
16772	-0.000	1.000	-0.000	...	2
0.0	1				
16773	-0.000	1.000	-0.000	...	2
0.0	0				
16774	-0.000	1.000	-0.000	...	2
0.0	0				
16775	-0.000	1.000	-0.000	...	2
0.0	0				
16776	-0.000	1.000	-0.000	...	2
0.0	0				


```
[16777 rows x 25 columns]
```

```
#outliers
```

```
import seaborn as sns
```

```
sns.boxplot(x=GT['casualties'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6a0f6ff890>
```



```
# calculate interquartile range
```

```
q25, q75 = np.percentile(GT['casualties'], 25),
```

```
np.percentile(GT['casualties'], 75)
```

```
iqr = q75 - q25
```

```
print('Percentiles: 25th=%.3f, 75th=%.3f, IQR=%.3f' % (q25, q75, iqr))
```

```
# calculate the outlier cutoff
```

```
cut_off = iqr * 1.5
```

```
lower, upper = q25 - cut_off, q75 + cut_off
```

```
# identify outliers
```

```
for x in GT['casualties']:
```

```
    if(x<lower or x>upper):
```

```
        GT.drop(GT.loc[(GT.casualties)==x].index, inplace=True, axis=0)
```

```
Percentiles: 25th=0.000, 75th=5.000, IQR=5.000
```

```
GT.shape
```

```
(14640, 25)
```

```
GT.success.value_counts()
```

```
1    14226
```

```
0      414
```

```
Name: success, dtype: int64
```

```
target= GT.success
```

```
GT.drop('success', axis=1, inplace= True)
```

```
target
```

```
0      1
```

```
1      1
```

```
2      1
```

```
3      1
```

```
4      1
```

```
..
```

```
16772   1
```

```
16773   1
```

```
16774   1
```

```
16775   1
```

```
16776   1
```

```
Name: success, Length: 14640, dtype: int64
```

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(GT, target,  
test_size=0.2)
```

```
from sklearn.linear_model import LogisticRegression
```

```
# all parameters not specified are set to their defaults
```

```
lg = LogisticRegression()
```

```
lg.fit(x_train, y_train)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/  
_logistic.py:818: ConvergenceWarning: lbfgs failed to converge  
(status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
```

```
LogisticRegression()
```

```
y_pred_lg = lg.predict(x_test)
```

```
from sklearn.metrics import roc_curve, auc, confusion_matrix,  
classification_report, accuracy_score
```



```

score_lg = accuracy_score(y_pred_lg,y_test)
score_lg


0.9702868852459017

print("train score - " + str(lg.score(x_train, y_train)))
print("test score - " + str(lg.score(x_test, y_test)))

train score - 0.9720799180327869
test score - 0.9702868852459017

#Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm_lg = confusion_matrix(y_test,y_pred_lg)
sns.set(font_scale=1.4)
sns.heatmap(cm_lg, annot=True)
plt.show()

```



```

print(classification_report(y_test, y_pred_lg))

```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	87
1	0.97	1.00	0.98	2841
accuracy			0.97	2928
macro avg	0.49	0.50	0.49	2928
weighted avg	0.94	0.97	0.96	2928

```

/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.7/dist-packages/sklearn/metrics/_classification.py:1308: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))

```

FEATURE SELECTION

(1) Low variance filter

```

from sklearn.feature_selection import VarianceThreshold
var_thres=VarianceThreshold(threshold=0)
var_thres.fit(GT)

VarianceThreshold(threshold=0)

var_thres.get_support()

array([ True,  True,  True,  True,  True,  True,  True,  True,  True,
        True, False,  True,  True,  True,  True,  True,  True,  True,
        True,  True,  True,  True,  True,  True])

GT.columns[var_thres.get_support()]

Index(['iyear_sin', 'iyear_cos', 'imonth_sin', 'imonth_cos', 'nkill',
      'nwound',
      'casualties', 'extended', 'country_txt', 'region_txt', 'crit2',
      'crit3',
      'multiple', 'vicinity', 'specificity', 'suicide',
      'attacktype1_txt',
      'weaptype1_txt', 'targettype1_txt', 'gname', 'propextent_txt',
      'ishostkid', 'INT_ANY'],
      dtype='object')

constant_columns = [column for column in GT.columns
                    if column not in
GT.columns[var_thres.get_support()]]

print(len(constant_columns))

```

```
for feature in constant_columns:
    print(feature)
```

```
crit1
```

```
GT= GT.drop(constant_columns,axis=1)
```

(2) High correlation filter

```
GT1.corr()
```

	iyear_sin	iyear_cos	imonth_sin	...	nkill	nwound
casualties						
iyear_sin	1.000000	0.963642	0.007827	...	0.052817	0.017822
0.024091						
iyear_cos	0.963642	1.000000	0.005302	...	0.055201	0.024010
0.029794						
imonth_sin	0.007827	0.005302	1.000000	...	-0.016187	-0.015842
-0.016282						
imonth_cos	-0.013271	-0.018387	0.021238	...	-0.009149	-0.002868
-0.003985						
nkill	0.052817	0.055201	-0.016187	...	1.000000	0.828968
0.877537						
nwound	0.017822	0.024010	-0.015842	...	0.828968	1.000000
0.995638						
casualties	0.024091	0.029794	-0.016282	...	0.877537	0.995638
1.000000						

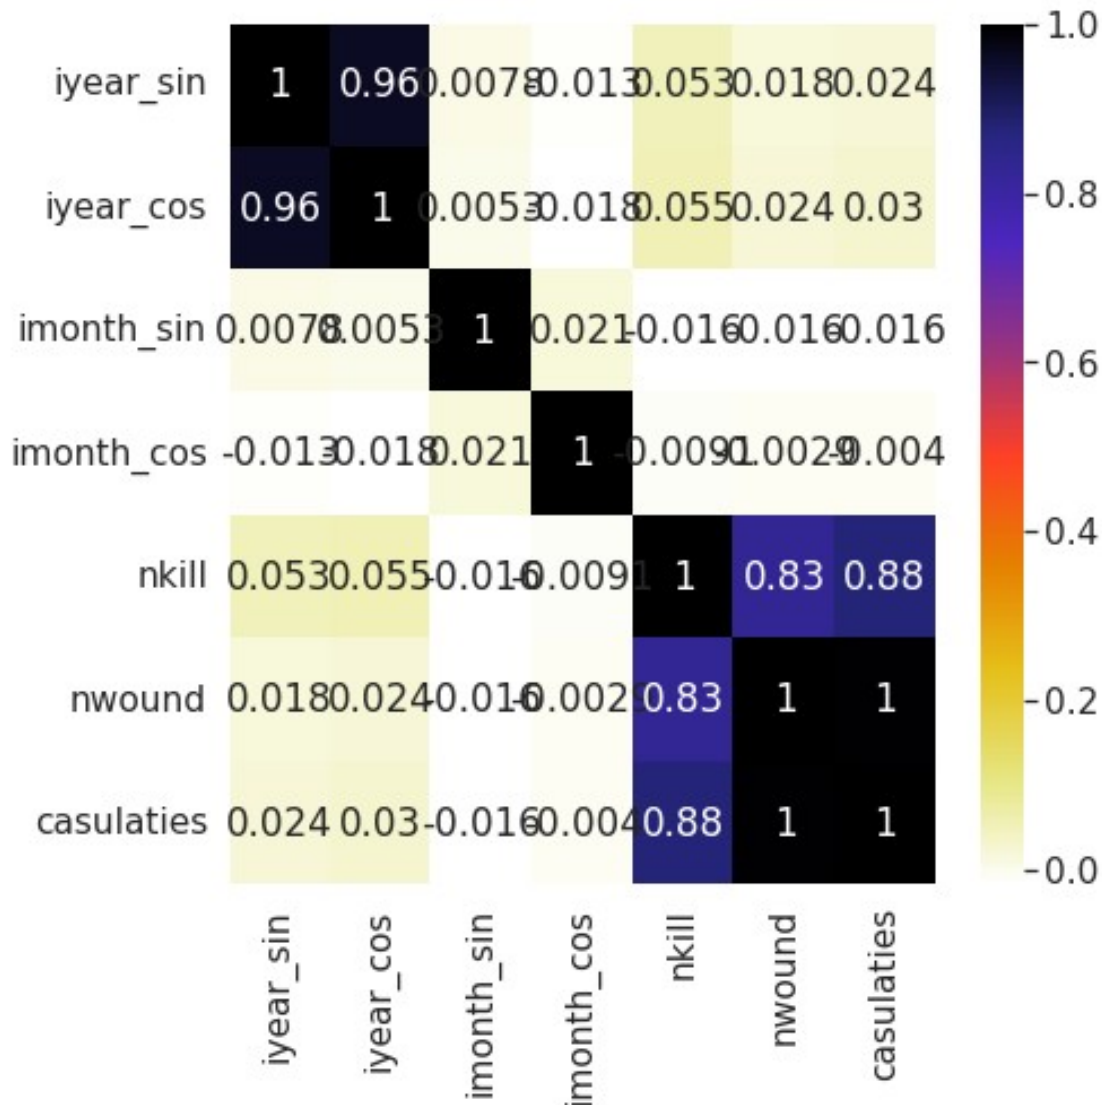
```
[7 rows x 7 columns]
```

```
plt.figure(figsize=(7,7))
```

```
cor =GT1.corr()
```

```
sns.heatmap(cor, annot=True, cmap=plt.cm.CMRmap_r)
```

```
plt.show()
```



*# with the following function we can select highly correlated features
it will remove the first feature that is correlated with anything
other feature*

```
def correlation(dataset, threshold):
    col_corr = [] # Set of all the names of correlated columns
    corr_matrix = dataset.corr()
    for i in range(len(corr_matrix.columns)):
        for j in range(i):
            if abs(corr_matrix.iloc[i, j]) > threshold: # we are
interested in absolute coeff value
                colname = corr_matrix.columns[i] # getting the name
of column
                col_corr.append([colname, corr_matrix.columns[j]])
    return col_corr
```

```
corr_features = correlation(GT1, 0.8)
len(corr_features)
```

```
4
```

```
corr_features
```

```
[['iyear_cos', 'iyear_sin'],
 ['nwound', 'nkill'],
 ['casualties', 'nkill'],
 ['casualties', 'nwound']]
```

```
GT= GT.drop(['nwound', 'nkill'], axis=1)
```

```
(3) Chi square test
```

```
GT2= GT[['extended', 'country_txt', 'region_txt', 'crit2', 'crit3',
'multiple', 'vicinity', 'specificity', 'suicide',
'attacktype1_txt', 'weaptype1_txt', 'targettype1_txt', 'gname',
'propextent_txt', 'ishostkid', 'INT_ANY']]
```

```
GT2
```

	extended	country_txt	region_txt	...	propextent_txt
ishostkid	INT_ANY				
0	0	123	6	...	2
0.0	0				
1	0	123	6	...	2
0.0	0				
2	0	124	7	...	2
0.0	0				
3	0	123	6	...	2
0.0	1				
4	0	123	6	...	2
0.0	1				
...
..	...				
16772	0	61	10	...	2
0.0	1				
16773	0	49	8	...	2
0.0	0				
16774	0	0	8	...	2
0.0	0				
16775	0	92	9	...	2
0.0	0				
16776	0	92	9	...	2
0.0	0				

```
[14640 rows x 16 columns]
```

```
target.shape
```

```

(14640,)

from sklearn.feature_selection import chi2
f_p_values= chi2(GT2,target)

f_p_values

(array([4.62302686e+00, 4.92754250e+01, 1.14699950e+01, 7.95235548e-
06,
        1.98795306e-06, 6.04375435e+01, 4.86155311e-01,
1.72947445e+00,
        1.74338265e+01, 7.97739094e+01, 1.99010202e+01,
9.87249583e+01,
        2.66908071e-01, 1.02792188e-01, 1.17250588e+01, 5.25044448e-
01]),
array([3.15455245e-02, 2.22431640e-12, 7.07288693e-04, 9.97749975e-
01,
        9.98875025e-01, 7.59504716e-15, 4.85647363e-01, 1.88478030e-
01,
        2.97484607e-05, 4.19799716e-19, 8.15567928e-06, 2.90122519e-
23,
        6.05413446e-01, 7.48504284e-01, 6.16640999e-04, 4.68697760e-
01]))

p_values=pd.Series(f_p_values[1])
p_values.index=GT2.columns
p_values

extended          3.154552e-02
country_txt       2.224316e-12
region_txt        7.072887e-04
crit2             9.977500e-01
crit3             9.988750e-01
multiple          7.595047e-15
vicinity          4.856474e-01
specificity       1.884780e-01
suicide           2.974846e-05
attacktype1_txt   4.197997e-19
weaptype1_txt     8.155679e-06
targetype1_txt    2.901225e-23
gname             6.054134e-01
propextent_txt    7.485043e-01
ishostkid         6.166410e-04
INT_ANY           4.686978e-01
dtype: float64

p_values.sort_index(ascending=False)

weaptype1_txt     8.155679e-06
vicinity          4.856474e-01
targetype1_txt    2.901225e-23
suicide           2.974846e-05

```

```
specificity          1.884780e-01
region_txt           7.072887e-04
propextent_txt       7.485043e-01
multiple             7.595047e-15
ishostkid            6.166410e-04
gname                6.054134e-01
extended             3.154552e-02
crit3                9.988750e-01
crit2                9.977500e-01
country_txt          2.224316e-12
attacktype1_txt      4.197997e-19
INT_ANY              4.686978e-01
dtype: float64
```

```
GT.drop(['INT_ANY', 'attacktype1_txt', 'country_txt', 'crit2', 'crit3', 'extended'], inplace=True, axis=1)
```

```
GT['success'] = target
```

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test = train_test_split(GT, target,
                                                    test_size=0.2)
```

```
from sklearn.linear_model import LogisticRegression
# all parameters not specified are set to their defaults
lg = LogisticRegression()
lg.fit(x_train, y_train)
```

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,
```

```
LogisticRegression()
```

```
y_pred_lg = lg.predict(x_test)
```

```
from sklearn.metrics import roc_curve, auc, confusion_matrix,
classification_report, accuracy_score
score_lg = accuracy_score(y_pred_lg, y_test)
score_lg
```


```
1.0
```

```
print("train score - " + str(lg.score(x_train, y_train)))
print("test score - " + str(lg.score(x_test, y_test)))
```

```
train score - 1.0
test score - 1.0
```

#Making the Confusion Matrix

```
from sklearn.metrics import confusion_matrix
cm_lg = confusion_matrix(y_test, y_pred_lg)
sns.set(font_scale=1.4)
sns.heatmap(cm_lg, annot=True)
plt.show()
```



```
print(classification_report(y_test, y_pred_lg))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	92
1	1.00	1.00	1.00	2836
accuracy			1.00	2928
macro avg	1.00	1.00	1.00	2928
weighted avg	1.00	1.00	1.00	2928

Smote

```
pip install imbalanced-learn
```



```

Requirement already satisfied: imbalanced-learn in
/usr/local/lib/python3.7/dist-packages (0.8.1)
Requirement already satisfied: scipy>=0.19.1 in
/usr/local/lib/python3.7/dist-packages (from imbalanced-learn) (1.4.1)
Requirement already satisfied: numpy>=1.13.3 in
/usr/local/lib/python3.7/dist-packages (from imbalanced-learn)
(1.19.5)
Requirement already satisfied: scikit-learn>=0.24 in
/usr/local/lib/python3.7/dist-packages (from imbalanced-learn) (1.0.1)
Requirement already satisfied: joblib>=0.11 in
/usr/local/lib/python3.7/dist-packages (from imbalanced-learn) (1.1.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.7/dist-packages (from scikit-learn>=0.24-
>imbalanced-learn) (3.0.0)

```

GT

	iyear_sin	iyear_cos	imonth_sin	...	propextent_txt
ishostkid	success				
0	-0.146	0.989	0.500	...	2
0.0	1				
1	-0.146	0.989	0.500	...	2
0.0	1				
2	-0.146	0.989	0.500	...	2
0.0	1				
3	-0.146	0.989	0.866	...	2
0.0	1				
4	-0.146	0.989	0.866	...	2
0.0	1				
...
.
16772	-0.000	1.000	-0.000	...	2
0.0	1				
16773	-0.000	1.000	-0.000	...	2
0.0	1				
16774	-0.000	1.000	-0.000	...	2
0.0	1				
16775	-0.000	1.000	-0.000	...	2
0.0	1				
16776	-0.000	1.000	-0.000	...	2
0.0	1				

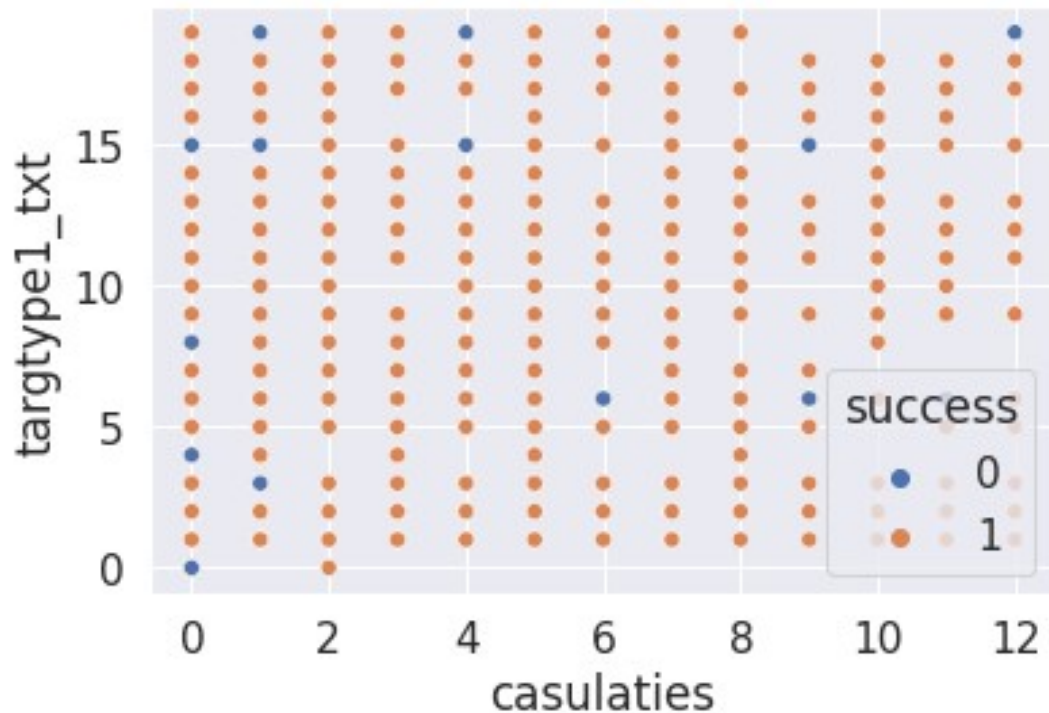
[14640 rows x 16 columns]

```

sns.scatterplot(data= GT, x= GT.casualties,
                 y= GT.targetype1_txt,
                 hue= GT.success)

```

<matplotlib.axes._subplots.AxesSubplot at 0x7f6a09e54c50>



```
# Oversample and plot imbalanced dataset with SMOTE
from collections import Counter
from sklearn.datasets import make_classification
from imblearn.over_sampling import SMOTE
from numpy import where

print("Before OverSampling, counts of label '1':
{}".format(sum(y_train == 1)))
print("Before OverSampling, counts of label '0': {} \
n".format(sum(y_train == 0)))

# import SMOTE module from imblearn library
# pip install imblearn (if you don't have imblearn in your system)
from imblearn.over_sampling import SMOTE
sm = SMOTE(random_state = 2)
x_train_res, y_train_res = sm.fit_resample(x_train, y_train.ravel())

print('After OverSampling, the shape of train_X:
{}'.format(x_train_res.shape))
print('After OverSampling, the shape of train_y: {} \
n'.format(y_train_res.shape))

print("After OverSampling, counts of label '1':
{}".format(sum(y_train_res == 1)))
print("After OverSampling, counts of label '0':
{}".format(sum(y_train_res == 0)))
```

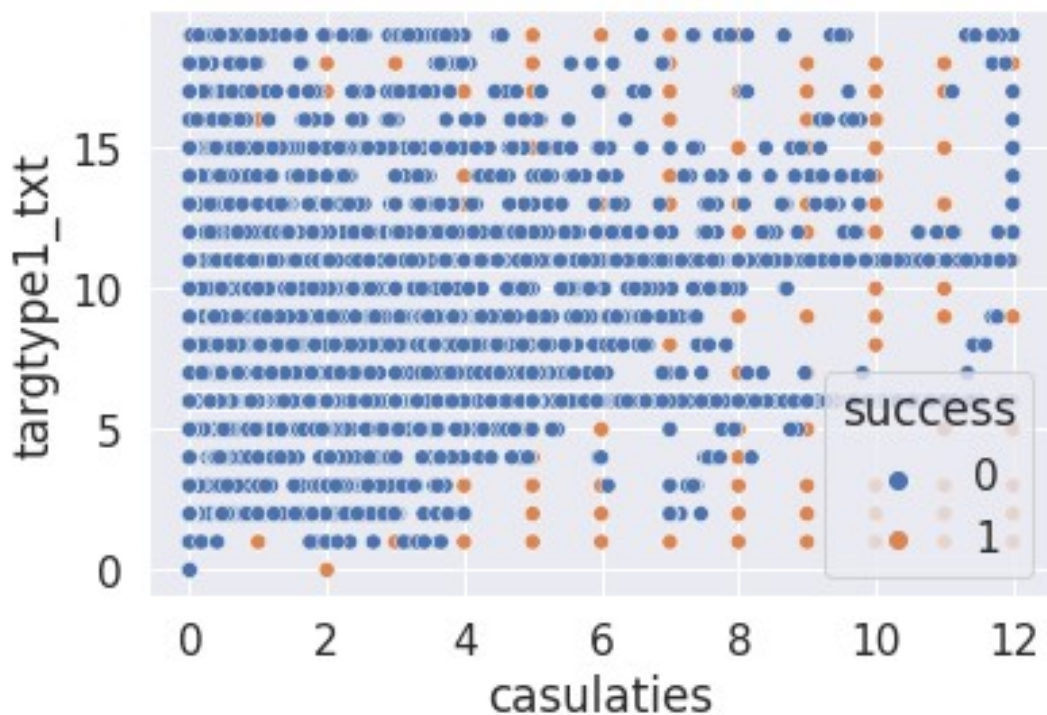
```
Before OverSampling, counts of label '1': 11390
Before OverSampling, counts of label '0': 322
```

```
After OverSampling, the shape of train_X: (22780, 16)
After OverSampling, the shape of train_y: (22780,)
```

```
After OverSampling, counts of label '1': 11390
After OverSampling, counts of label '0': 11390
```

```
sns.scatterplot(x= x_train_res.casualties,
                y= x_train_res.target1_txt,
                hue= x_train_res.success)
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f6a0f94e510>
```



```
lr1 = LogisticRegression()
lr1.fit(x_train_res, y_train_res.ravel())
predictions = lr1.predict(x_test)
```

```
# print classification report
print(classification_report(y_test, predictions))
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	92
1	1.00	1.00	1.00	2836
accuracy			1.00	2928

macro avg	1.00	1.00	1.00	2928
weighted avg	1.00	1.00	1.00	2928

```
/usr/local/lib/python3.7/dist-packages/sklearn/linear_model/_logistic.py:818: ConvergenceWarning: lbfgs failed to converge
(status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (`max_iter`) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>
Please also refer to the documentation for alternative solver options:


https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression
`extra_warning_msg=_LOGISTIC_SOLVER_CONVERGENCE_MSG,`

TIME SERIES ANALYSIS - ANNUAL DATA

Time series data is a collection of quantities that are assembled over **even** intervals in time and ordered chronologically

```
GT['Date'].min(), GT['Date'].max()
```

```
(Timestamp('1970-01-02 00:00:00'), Timestamp('2017-12-31 00:00:00'))
```



```
df = GT.groupby("Date")["ID"].count().reset_index()
df['year'] = [d.year for d in df.Date]
df['month'] = [d.strftime('%b') for d in df.Date]
df
```

	Date	ID	year	month
0	1970-01-02	1	1970	Jan
1	1970-01-03	1	1970	Jan
2	1970-01-15	1	1970	Jan
3	1970-02-08	2	1970	Feb
4	1970-02-13	1	1970	Feb
...
6373	2017-12-27	2	2017	Dec
6374	2017-12-28	1	2017	Dec

```
6375 2017-12-29    1  2017    Dec
6376 2017-12-30    2  2017    Dec
6377 2017-12-31    4  2017    Dec
```


```
[6378 rows x 4 columns]
```

```
fig, axes = plt.subplots(1, 2, figsize=(20,7), dpi= 80)
sns.boxplot(x='year', y='ID', data=df, ax=axes[0])
sns.boxplot(x='month', y='ID', data=df)
```

```
# Set Titles
```

```
axes[0].set_title('Year-wise Box Plot\n(The Trend)', fontsize=18);
axes[0].tick_params(labelrotation=45)
```

```
axes[1].set_title('Month-wise Box Plot\n(The Seasonality)',
fontsize=18)
plt.show()
```




```
fig, axes = plt.subplots(1, 2, figsize=(20,7), dpi= 80)
sns.boxplot(x='year', y='ID', data=df[3500:5700], ax=axes[0])
sns.boxplot(x='month', y='ID', data=df[3500:5700])
```

```
# Set Titles
```

```
axes[0].set_title('Year-wise Box Plot', fontsize=18);
axes[0].tick_params(labelrotation=45)
```

```
axes[1].set_title('Month-wise Box Plot', fontsize=18)
plt.show()
```



Date wise grouping does not have even interval

=>Year wise grouping

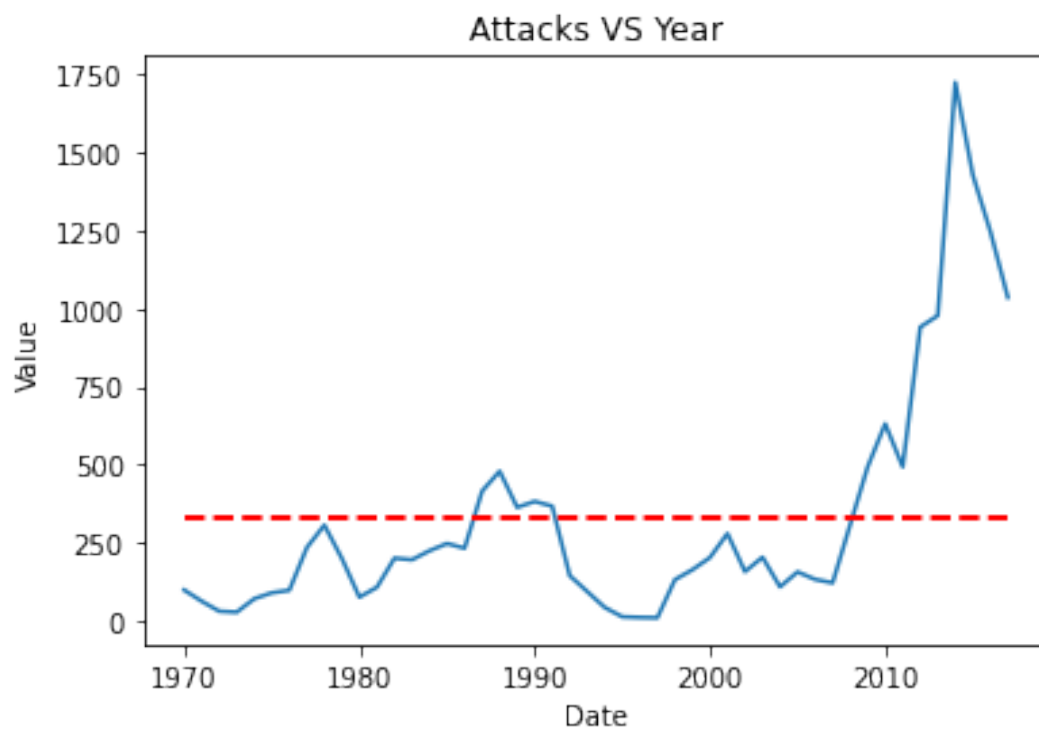

Year - Number of attacks - Casualties

```
GT1 = GT.groupby("Year")["ID"].count().reset_index()
GT1["Casualties"] =
GT.groupby('Year').Casualties.sum().reset_index()["Casualties"]
GT1.head()
```

	Year	ID	Casualties
0	1970	98	155.0
1	1971	62	9.0
2	1972	30	144.0
3	1973	27	161.0
4	1974	70	671.0

```
def plot_df(df, x, y, title="", xlabel='Date', ylabel='Value',
dpi=100):
    fig, ax = plt.subplots()
    ax.plot(x, y)
    y_avg = [np.mean(y)] * len(y)
    ax.plot(x, y_avg, color='red', lw=2, ls='--')
    plt.gca().set(title=title, xlabel=xlabel, ylabel=ylabel)
    plt.show()
```


```
plot_df(GT1, x=GT1.Year, y=GT1.Casualties, title='Casualties VS
Year')
print()
plot_df(GT1, x=GT1.Year, y=GT1.ID, title='Attacks VS Year')
```



Stationary or Non-stationary ?

Stationary if its **mean and variance are time invariant**

- A stationary time series will be mean reverting in nature, i.e. it will tend to return to its mean and fluctuations around the mean will have roughly equal amplitudes.
- A stationary time series will not drift too far away from its mean because of its finite constant variance.
- A non-stationary time series, on the contrary, will have a time varying variance or a time varying mean or both, and will not tend to revert back to its mean.



Red - Stationary

Blue - Non-Stationary

#Making Sure all the years are included

```
dd = GT1[["Year","ID"]]
all_year = pd.DataFrame({'Year':list(range(1970,2018))})
```

#Left join your main data on dates data

```
dd = pd.merge(all_year, dd, on='Year', how="left")
dd.fillna(0, inplace=True)
dd = dd.set_index("Year")
dd
```

	ID
Year	
1970	98.0
1971	62.0
1972	30.0
1973	27.0
1974	70.0
1975	89.0
1976	97.0
1977	234.0

1978	306.0
1979	199.0
1980	75.0
1981	107.0
1982	200.0
1983	195.0
1984	224.0
1985	247.0
1986	232.0
1987	415.0
1988	479.0
1989	363.0
1990	382.0
1991	367.0
1992	144.0
1993	0.0
1994	42.0
1995	12.0
1996	10.0
1997	9.0
1998	131.0
1999	163.0
2000	202.0
2001	279.0
2002	157.0
2003	203.0
2004	108.0
2005	156.0
2006	132.0
2007	121.0
2008	307.0
2009	493.0
2010	631.0
2011	493.0
2012	941.0
2013	978.0
2014	1726.0
2015	1429.0
2016	1248.0
2017	1036.0

How to test for stationarity?

1. *Plotting Rolling Statistics*
2. *Augmented Dickey Fuller test (ADH Test)*

This test will generate critical values and a p-value, which will allow us to accept or reject the null hypothesis that there is no stationarity. (*Null hypothesis: Non Stationary*)


- Reject -> Stationary

- Accept -> Non Stationary

Plotting Rolling Statistics

```
def stat_plot(df):
    rolling_mean = df.rolling(7).mean()
    rolling_std = df.rolling(7).std()
    plt.plot(df, color="blue", label="Original Data")
    plt.plot(rolling_mean, color="red", label="Rolling Mean Attacks")
    plt.plot(rolling_std, color="black", label = "Rolling Standard
Deviation in Attacks")
    plt.title("Attacks-Time Series, Rolling Mean, Standard Deviation")
    plt.legend(loc="best")
```

stat_plot(dd)



Augmented Dickey Fuller test (ADF Test)

```
# ADF Test
def adf(df):
    adft = adfuller(df['ID'], autolag='AIC')
    output_df = pd.DataFrame({"Values":
[adft[0],adft[1],adft[2],adft[3], adft[4]['1%'], adft[4]['5%'],
adft[4]['10%']] , "Metric":["Test Statistics","p-value","No. of lags
used","Number of observations used",
"critical
value (1%)", "critical value (5%)", "critical value (10%)"]})
    print(output_df)

adf(dd)
```

	Values	Metric
0	-0.946668	Test Statistics
1	0.772216	p-value
2	0.000000	No. of lags used
3	47.000000	Number of observations used
4	-3.577848	critical value (1%)
5	-2.925338	critical value (5%)
6	-2.600774	critical value (10%)

p-value(0.772216) > 0.05

=>Accept Null Hypothesis

=> Non Stationary

Autocorrelation

- This is a measure of how correlated time series data is at a given point in time with past values, which has huge implications across many industries.
- For example, if our GT data has strong autocorrelation, we can assume that high attack numbers today suggest a strong likelihood that they will be high tomorrow as well.

```
autocorrelation_lag1 = dd['ID'].autocorr(lag=1)
print("Lag1: ", autocorrelation_lag1)
```

```
autocorrelation_lag3 = dd['ID'].autocorr(lag=3)
print("Lag3: ", autocorrelation_lag3)
```


```
autocorrelation_lag6 = dd['ID'].autocorr(lag=6)
print("Lag6: ", autocorrelation_lag6)
```

```
Lag1:  0.9065891322854303
Lag3:  0.6947269909158585
Lag6:  0.35728230161986707
```

=> Corelation decreases long term

Decomposition

```
decompose = seasonal_decompose(dd['ID'],model='additive', period=7)
decompose.plot()
plt.show()
```




Non-Stationary to Stationary

Making the Difference

```
ddd = dd.diff().dropna()
adf(ddd)
plot_df(ddd, x=ddd.index, y=ddd.ID, title='Attacks VS Year')
```

	Values	Metric
0	-4.089281	Test Statistics
1	0.001009	p-value
2	2.000000	No. of lags used
3	44.000000	Number of observations used
4	-3.588573	critical value (1%)
5	-2.929886	critical value (5%)
6	-2.603185	critical value (10%)




p-value (0.001009) < 0.05

=> Reject Null Hypothesis

=> Stationary

stat_plot(ddd)



Log Transform - Moving average

```
dd_log = np.log(dd)
plt.plot(dd_log)
```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:1:
RuntimeWarning: divide by zero encountered in log
    """Entry point for launching an IPython kernel.
```

```
[<matplotlib.lines.Line2D at 0x7f21f15b98d0>]
```




```
moving_avg = dd_log.rolling(4).mean()
plt.plot(dd_log)
plt.plot(moving_avg, color='red')
[<matplotlib.lines.Line2D at 0x7f21f148bf50>]
```



```
dd_log_moving_avg_diff = dd_log - moving_avg
dd_log_moving_avg_diff.dropna(inplace=True)
```

```
adf(dd_log_moving_avg_diff)
stat_plot(dd_log_moving_avg_diff)
```

	Values	Metric
0	-4.581888	Test Statistics
1	0.000139	p-value
2	1.000000	No. of lags used
3	39.000000	Number of observations used
4	-3.610400	critical value (1%)
5	-2.939109	critical value (5%)
6	-2.608063	critical value (10%)



p-value (0.00139) < 0.05

=> Reject Null Hypothesis

=> Stationary

FORECASTING - ARIMA

Time series forecasting allows us to predict future values in a time series given current and past data

Forecast the number of attacks using ARIMA

auto_arima : Fit best ARIMA model to univariate time series (ensemble)

#Train-Test Split

```
def ttsplit(df):
    df['Year'] = df.index
```



```

train = df[df['Year'] < 2010]
train['train'] = train['ID']
del train['Year']
del train['ID']
test = df[df['Year'] >= 2010]
del test['Year']
test['test'] = test['ID']
del test['ID']
del df['Year']

plt.plot(train, color = "black")
plt.plot(test, color = "red")
plt.title("Train/Test split for Passenger Data")
plt.ylabel("Passenger Number")
plt.xlabel('Year-Month')
sns.set()
plt.show()

return train,test

```

#1

```
train, test = ttsplit(ddd)
```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:5:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```


See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:10:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

```

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
 # Remove the CWD from sys.path while we load stuff.



```

model = auto_arima(train, trace=True, error_action='ignore',
suppress_warnings=True)
#trace - If TRUE, the list of ARIMA models considered will be
reported.
model.fit(train)
forecast = model.predict(n_periods=len(test))
forecast = pd.DataFrame(forecast,index =
test.index,columns=['Prediction'])

```

Performing stepwise search to minimize aic

```

ARIMA(2,0,2)(0,0,0)[0] intercept : AIC=468.550, Time=0.30 sec
ARIMA(0,0,0)(0,0,0)[0] intercept : AIC=464.802, Time=0.01 sec
ARIMA(1,0,0)(0,0,0)[0] intercept : AIC=465.286, Time=0.06 sec
ARIMA(0,0,1)(0,0,0)[0] intercept : AIC=464.474, Time=0.04 sec
ARIMA(0,0,0)(0,0,0)[0]          : AIC=463.304, Time=0.01 sec
ARIMA(1,0,1)(0,0,0)[0] intercept : AIC=464.638, Time=0.11 sec

```

Best model: ARIMA(0,0,0)(0,0,0)[0]

Total fit time: 0.541 seconds

forecast

	Prediction
Year	
2010	0.0
2011	0.0
2012	0.0

```

2013      0.0
2014      0.0
2015      0.0
2016      0.0
2017      0.0

```

```

rms = sqrt(mean_squared_error(test,forecast))
print("RMSE: ", rms)

```

```


RMSE:  347.41527744185345

```

```

plt.plot(train, color = "black")
plt.plot(test, color = "red")
plt.plot(forecast, color = "green")
plt.title("Train/Test split for Passenger Data")
plt.ylabel("Passenger Number")
plt.xlabel('Year-Month')
sns.set()
plt.show()

```



#2

```

train, test = ttsplit(dd_log_moving_avg_diff)

```

```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:5:
SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead


```

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
"""  
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:10:  
SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame.  
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:
https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
# Remove the CWD from sys.path while we load stuff.
```



```
model = auto_arima(train, trace=True, error_action='ignore',  
suppress_warnings=True)  
#trace - If TRUE, the list of ARIMA models considered will be  
reported.  
model.fit(train)  
forecast = model.predict(n_periods=len(test))  
forecast = pd.DataFrame(forecast, index =  
test.index, columns=['Prediction'])
```

Performing stepwise search to minimize aic

```
ARIMA(2,0,2)(0,0,0)[0] intercept : AIC=inf, Time=0.26 sec  
ARIMA(0,0,0)(0,0,0)[0] intercept : AIC=61.696, Time=0.02 sec  
ARIMA(1,0,0)(0,0,0)[0] intercept : AIC=57.021, Time=0.03 sec
```

```
ARIMA(0,0,1)(0,0,0)[0] intercept : AIC=51.570, Time=0.05 sec
ARIMA(0,0,0)(0,0,0)[0]           : AIC=64.080, Time=0.02 sec
ARIMA(1,0,1)(0,0,0)[0] intercept : AIC=inf, Time=0.14 sec
ARIMA(0,0,2)(0,0,0)[0] intercept : AIC=inf, Time=0.10 sec
ARIMA(1,0,2)(0,0,0)[0] intercept : AIC=inf, Time=0.23 sec
ARIMA(0,0,1)(0,0,0)[0]           : AIC=inf, Time=0.06 sec
```

Best model: ARIMA(0,0,1)(0,0,0)[0] intercept
Total fit time: 0.927 seconds

forecast

	Prediction
Year	
2010	0.485146
2011	0.210597
2012	0.210597
2013	0.210597
2014	0.210597
2015	0.210597
2016	0.210597
2017	0.210597

```
rms = sqrt(mean_squared_error(test,forecast))
print("RMSE: ", rms)
```

RMSE: 0.26240347605662645

```
plt.plot(train, color = "black")
plt.plot(test, color = "red")
plt.plot(forecast, color = "green")
plt.title("Train/Test split for Passenger Data")
plt.ylabel("Passenger Number")
plt.xlabel('Year-Month')
sns.set()
plt.show()
```

